



mongoDB

**What?**



Humongous

Because it can store lots and lots of data

# Foundations of Document Databases with MongoDB

---

INTRODUCTION TO MONGODB

# Agenda



**What is NoSQL?**

**CAP Theorem**

**Types of NoSQL Databases**

**Advantages of Document Database**

**Why Use MongoDB?**

**JSON vs BSON**

**SQL Terms vs MongoDB Terms**

# What is NoSQL?

NoSQL = Not Only SQL

# JSON (BSON) Data Format

```
{  
    "name": "Max",  
    "age": 29,  
    "address":  
        {  
            "city": "Munich"  
        },  
    "hobbies": [  
        { "name": "Cooking" },  
        { "name": "Sports" }  
    ]  
}
```

# Relational SQL vs Document DB

MySQL	MongoDB
Structured Query Language (SQL)	MongoDB Query Language (MQL)
Predefined schema	Dynamic schema (JSON based)
Relational keys (foreign key)	No foreign key
Triggers	No triggers
ACID properties	CAP theorem
Vertically scalable	Horizontal scalable

# Three Important Questions



Is your data structured or unstructured?



What is your scalability strategy for infrastructure?



How comfortable are your devs with Object Relational Mapping?

# CAP Theorem

C

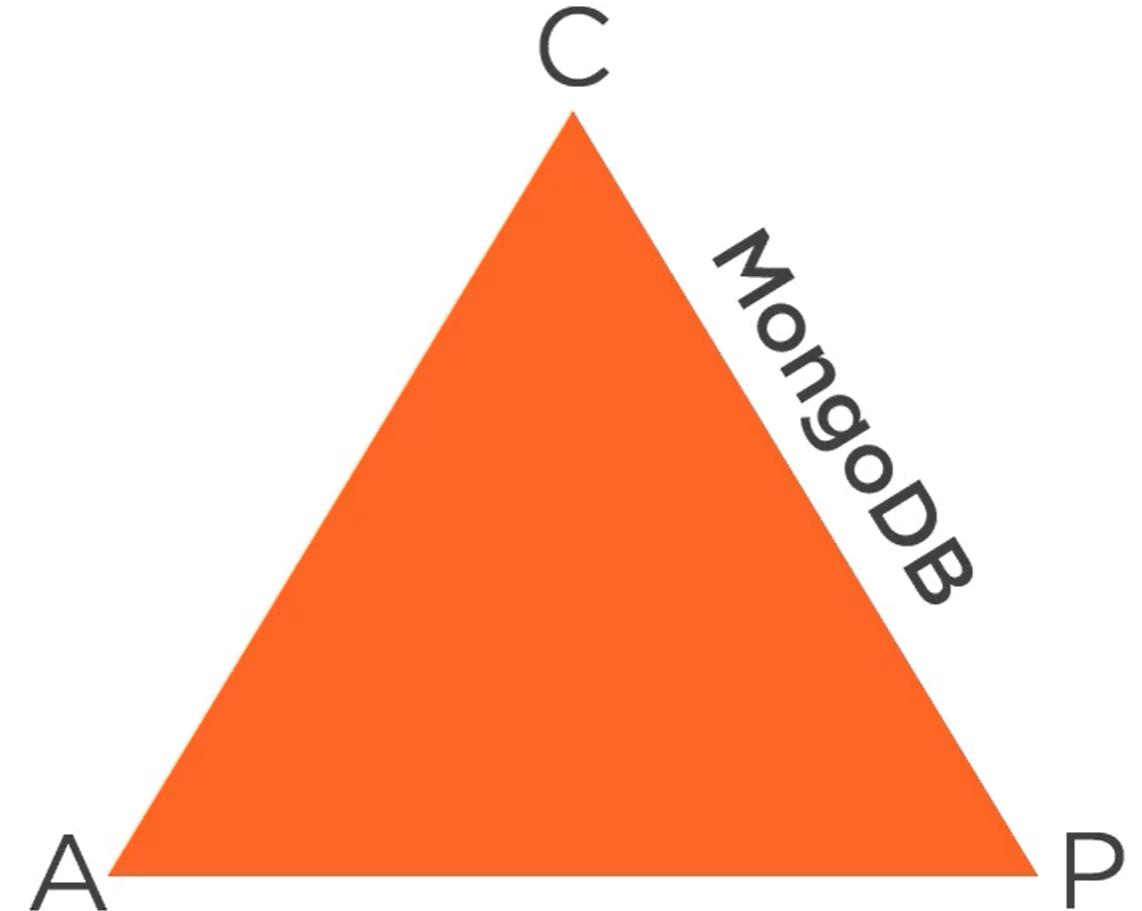
Consistency

A

Availability

P

Partition Tolerance



# Advantages of Document Database



**Intuitive Data Model**

**Flexible Schema**

**Universal JSON documents**

**Query Data Anyway**

**Distributed Scalable Database**

# Why Use MongoDB?



- Open Source and Free (Community Ed)**
- Document Database**
- High Performance**
- Rich Query Language**
- High Availability**
- Horizontal Scalability**
- Multiple Storage Engine**

# SQL Terms vs MongoDB Terms

SQL Terms	MongoDB Terms
Database	Database
Table	Collection
Row	JSON / BSON Document
Column	Field
Index	Index

# SQL Terms vs MongoDB Terms

SQL Terms	MongoDB Terms
Database	Database
Table	Collection
Row	Document
Column	Field
Index	Index
Table joins	\$lookup
Primary key	Primary Key
Transactions	Transactions

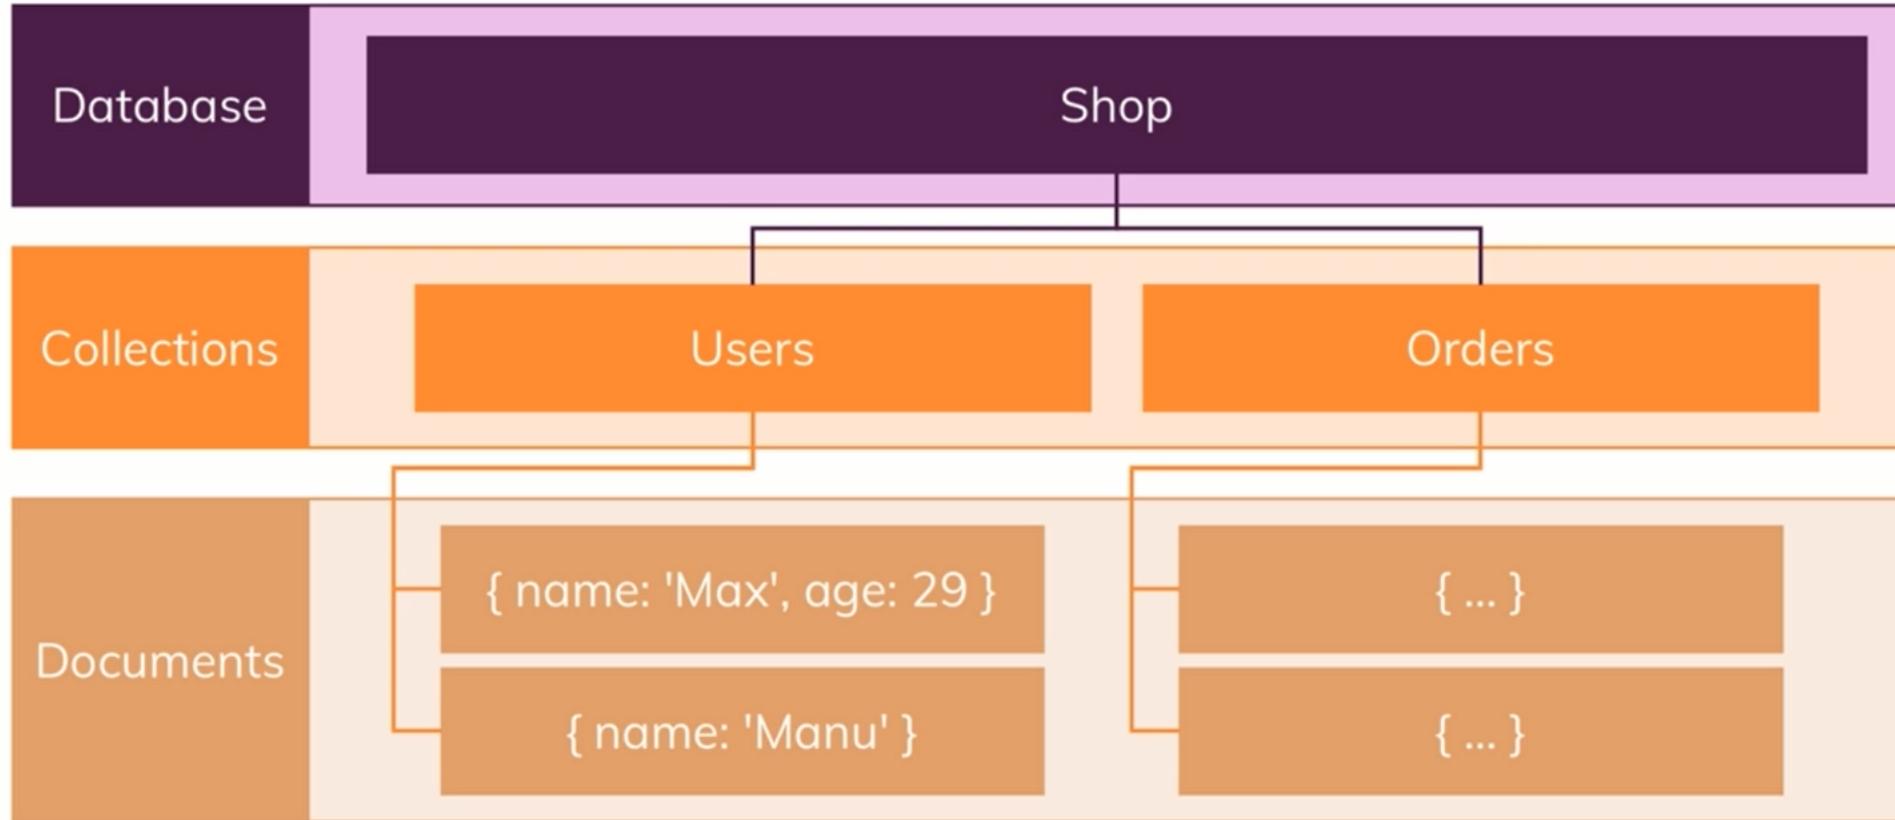
# JSON (BSON) Data Format

```
{  
    "name": "Max",  
    "age": 29,  
    "address":  
        {  
            "city": "Munich"  
        },  
    "hobbies": [  
        { "name": "Cooking" },  
        { "name": "Sports" }  
    ]  
}
```

# JSON vs BSON

	<b>JSON (JavaScript Object Notation)</b>	<b>BSON (Binary JSON)</b>
<b>Encoding</b>	UTF-8 String	Binary
<b>Data Support</b>	String, Boolean, Number, Array	String, Boolean, Number (Integer, Float, Long, Decimal128...), Array, Date, Raw Binary
<b>Readability</b>	Human and Machine	Machine Only

# How it works



# BSON Data Structure

No Schema!



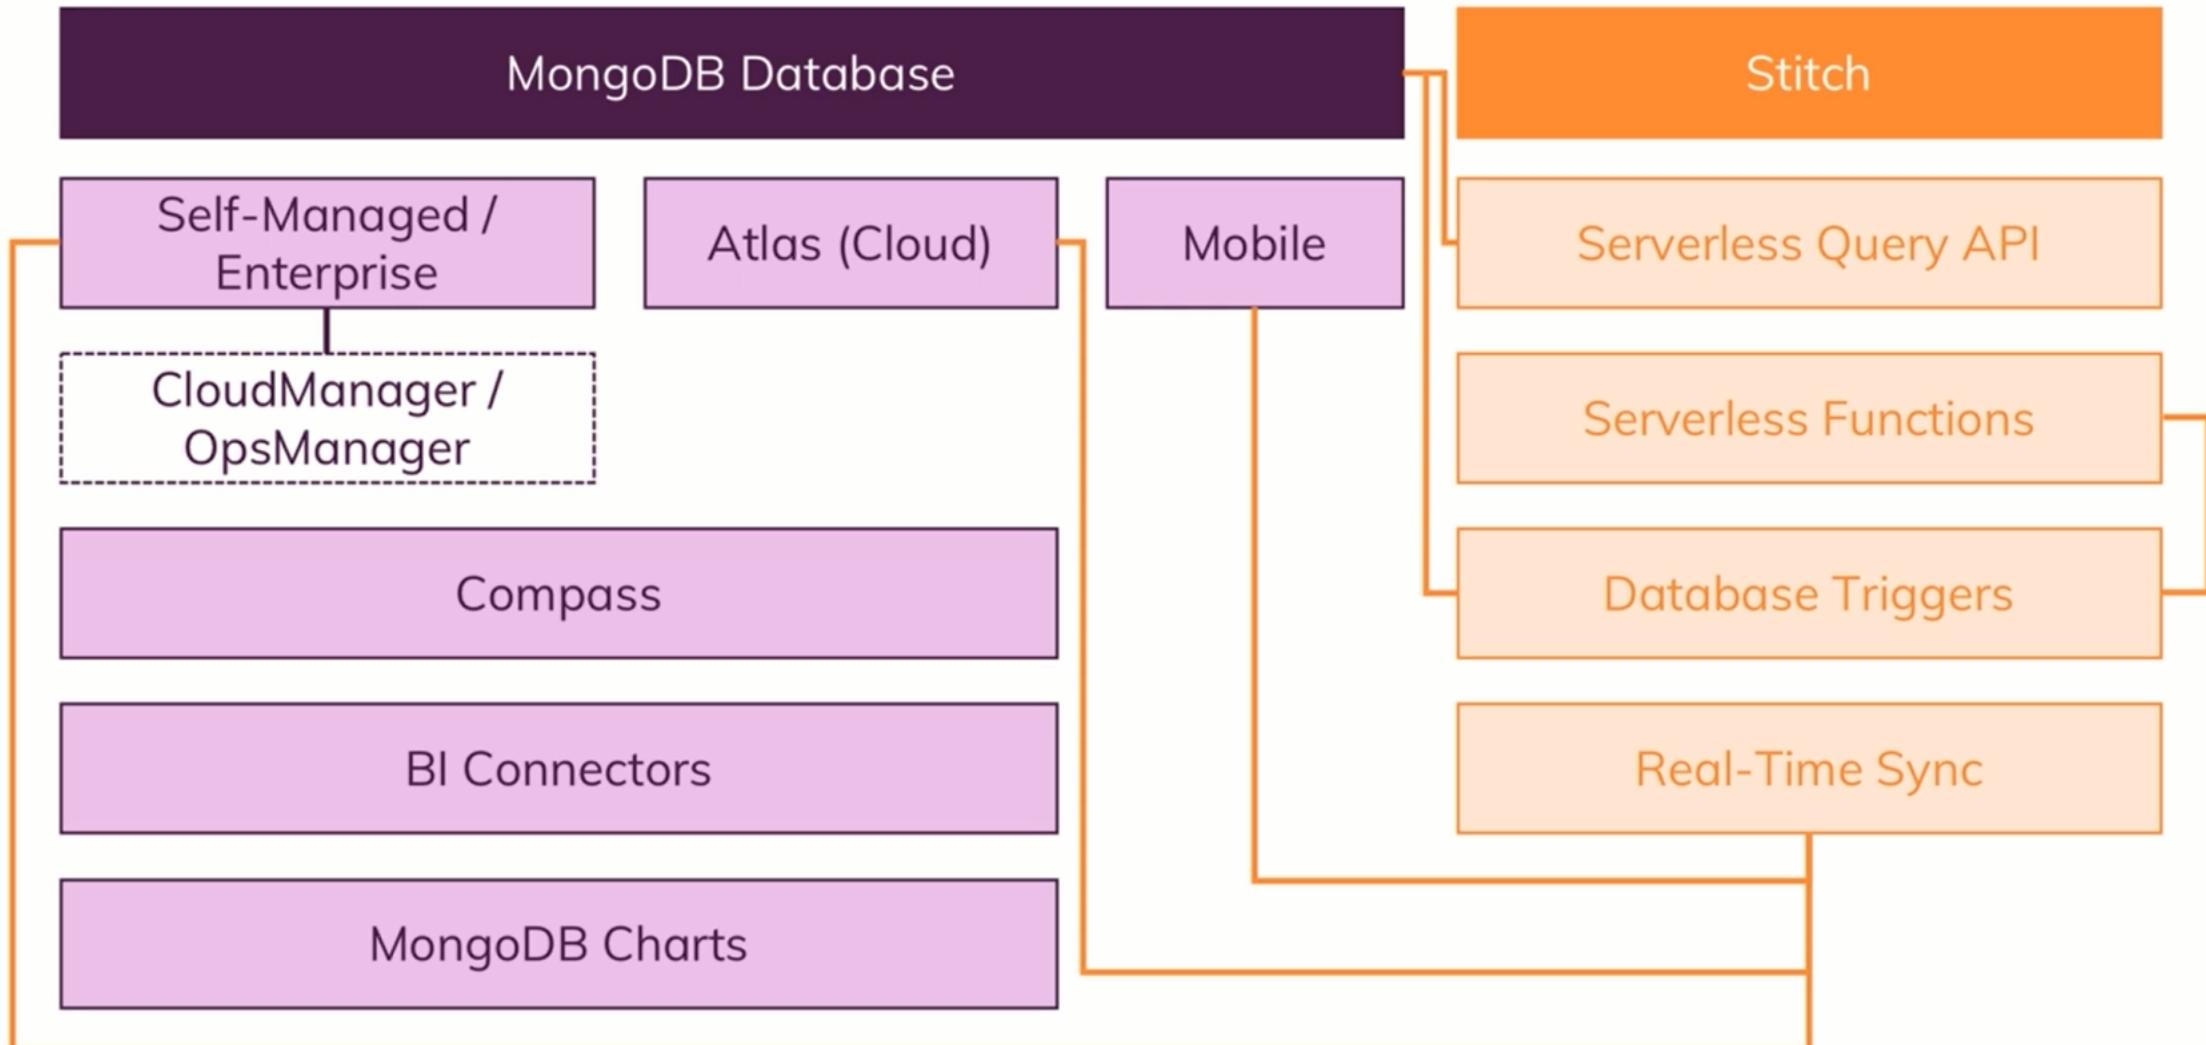
# Relations

No / Few Relations!

Relational Data needs to be merged **manually**

Kind of...

# MongoDB Ecosystem



“

# INSTALL MONGO DB

”

[HTTPS://WWW.MONGODB.COM/TRY/DOWNLOAD/COMMUNITY](https://www.mongodb.com/try/download/community)

MongoDB Community Serve

MongoDB offers both an Enterprise and Community version of its powerful distributed document database. The community version offers the flexible document model along with ad hoc queries, indexing, and real time aggregation to provide powerful ways to access and analyze your data. As a distributed system you can scale horizontally with native sharding.

The MongoDB Enterprise Server gives you all of this and more.



## Available Downloads

- Version 4.4.2 (current)
- Platform Windows
- Package msi

[Current releases & packages](#)  
[Development releases](#)  
[Archived releases](#)  
[Changelog](#)  
[Release Notes](#)

MongoDB Ops Manager

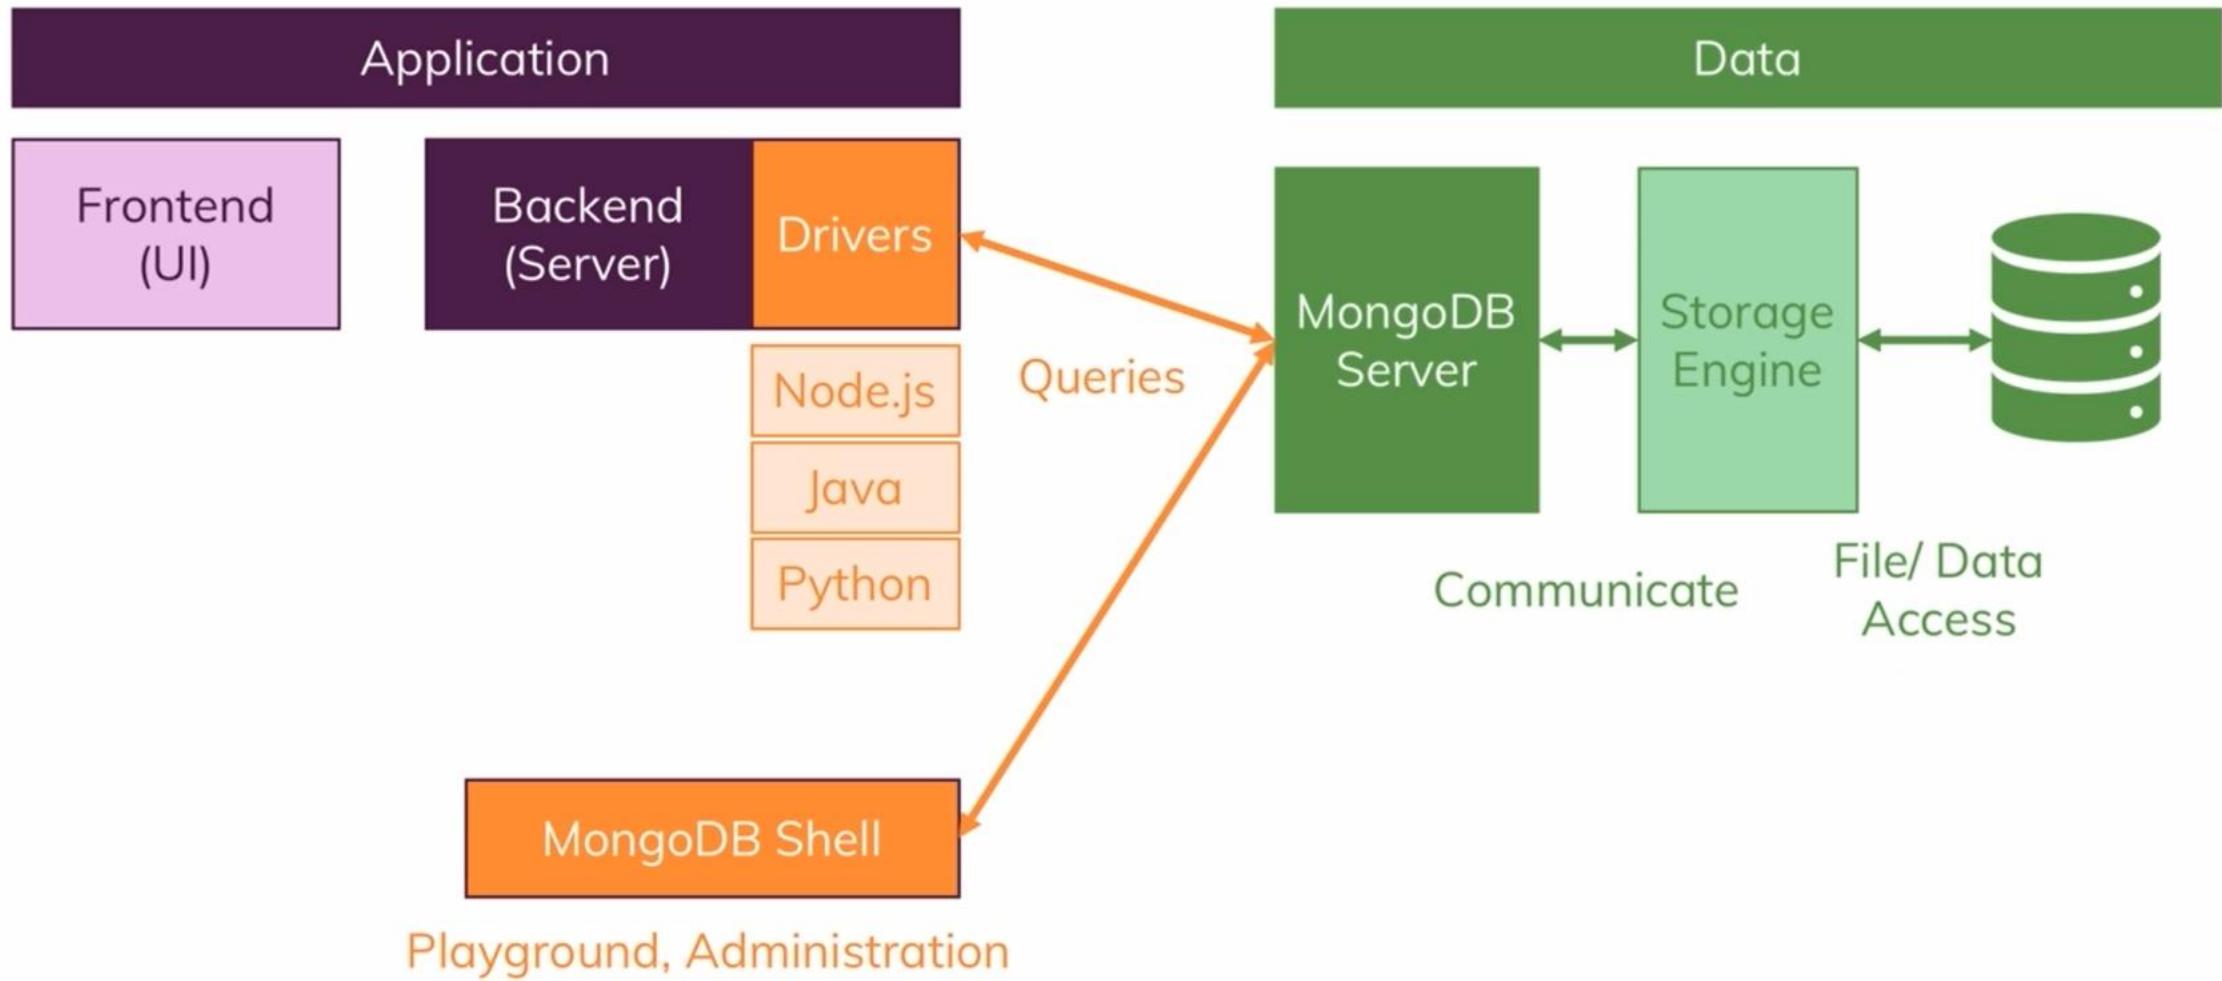
MongoDB Enterprise Kubernetes Operator

MongoDB Community Kubernetes Operator

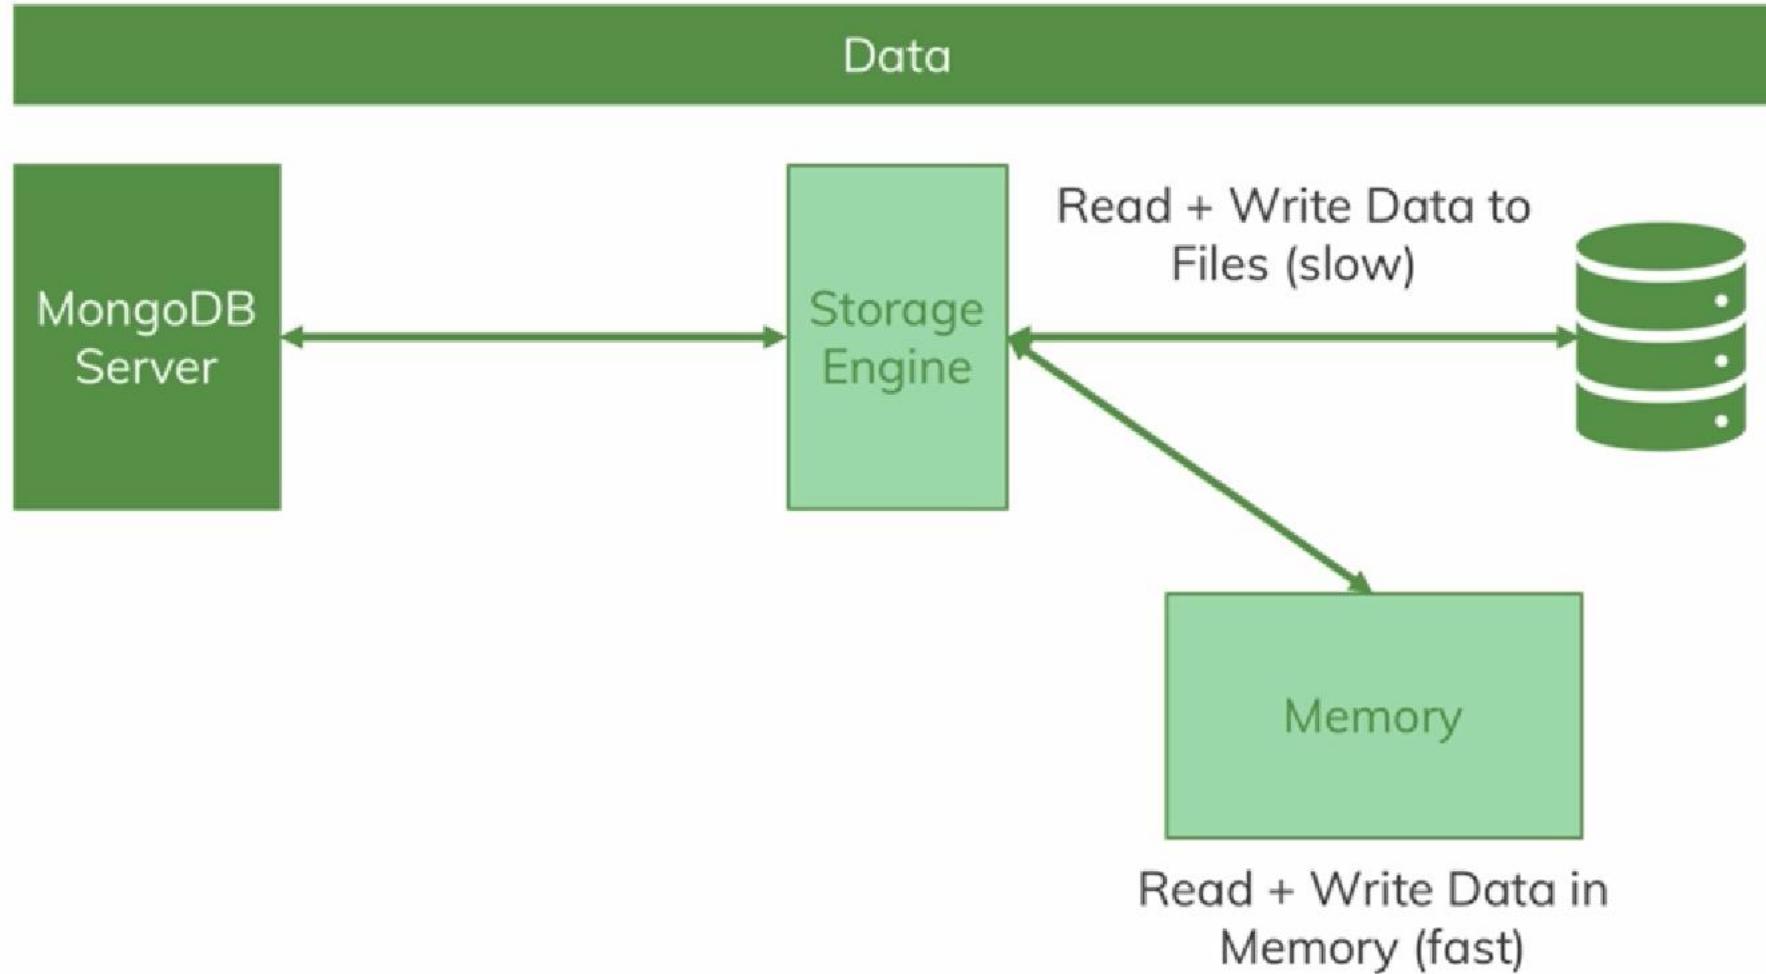
MongoDB Charts



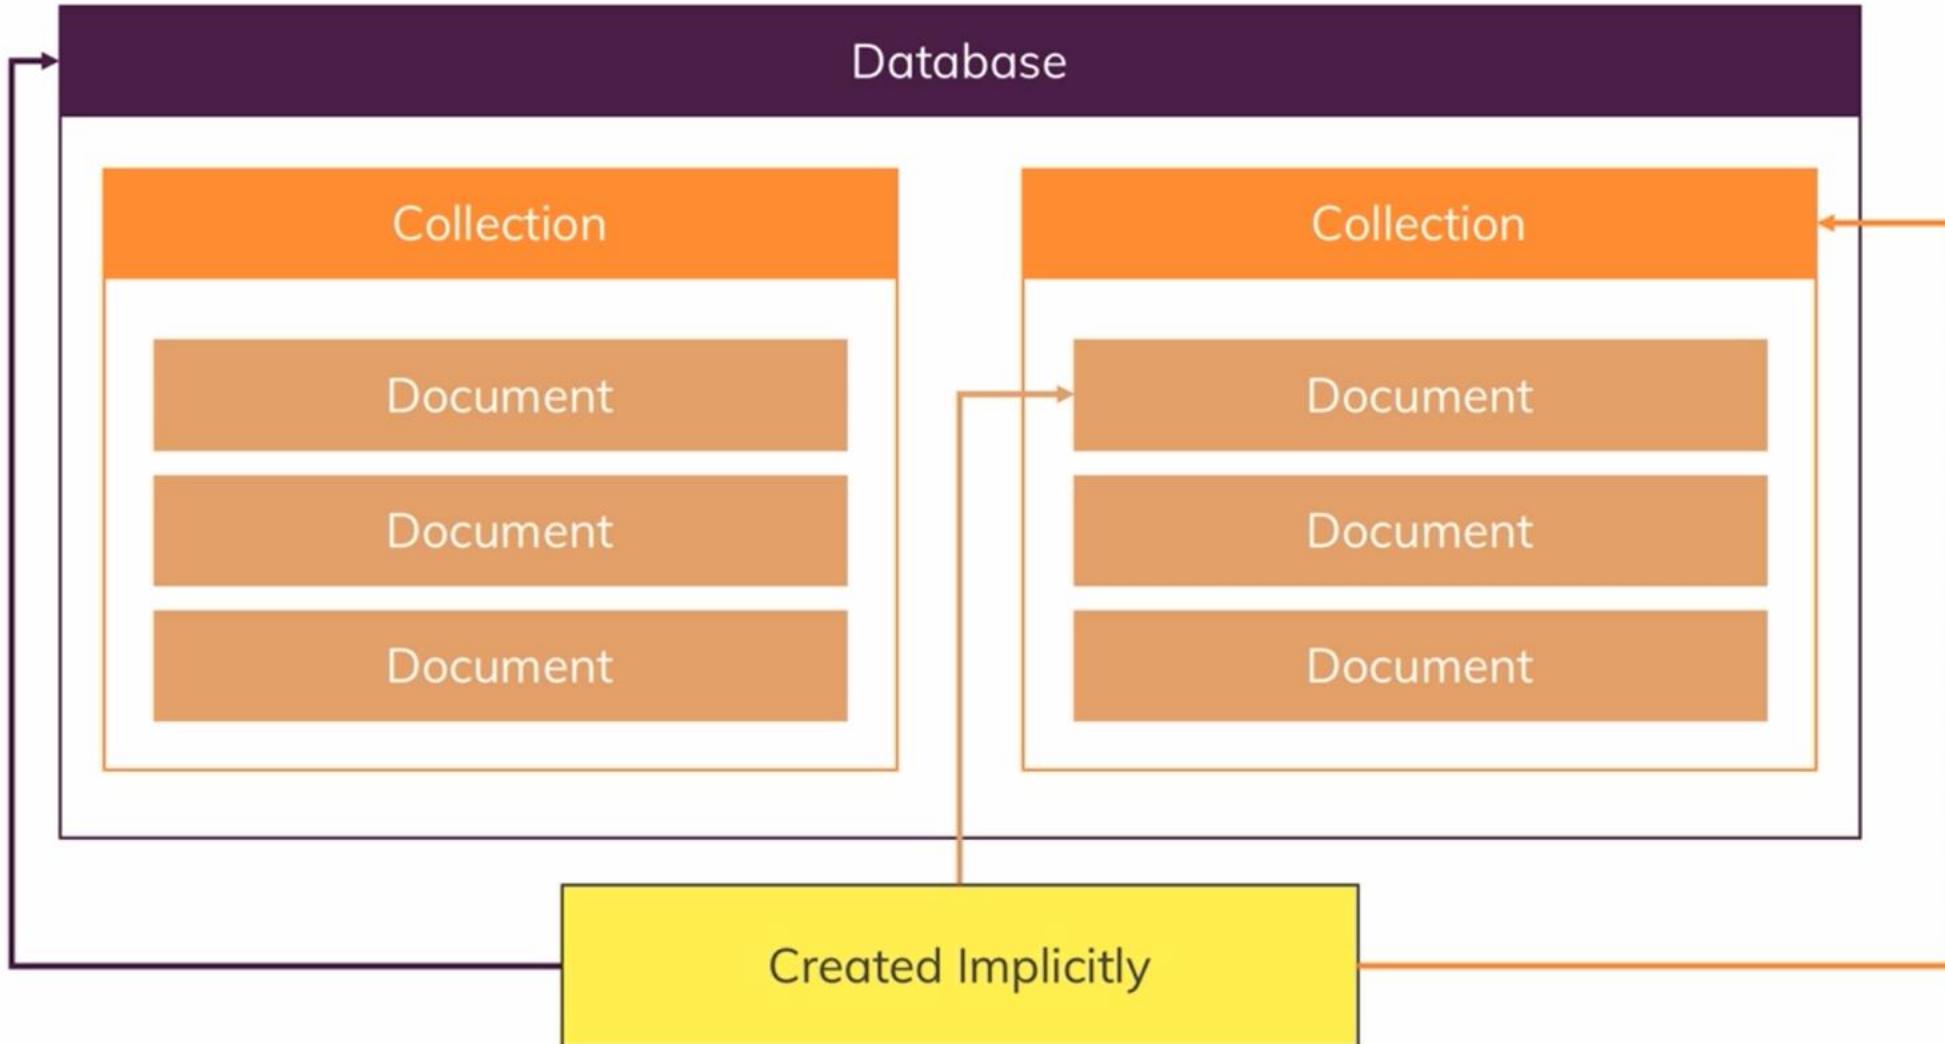
# Working with MongoDB



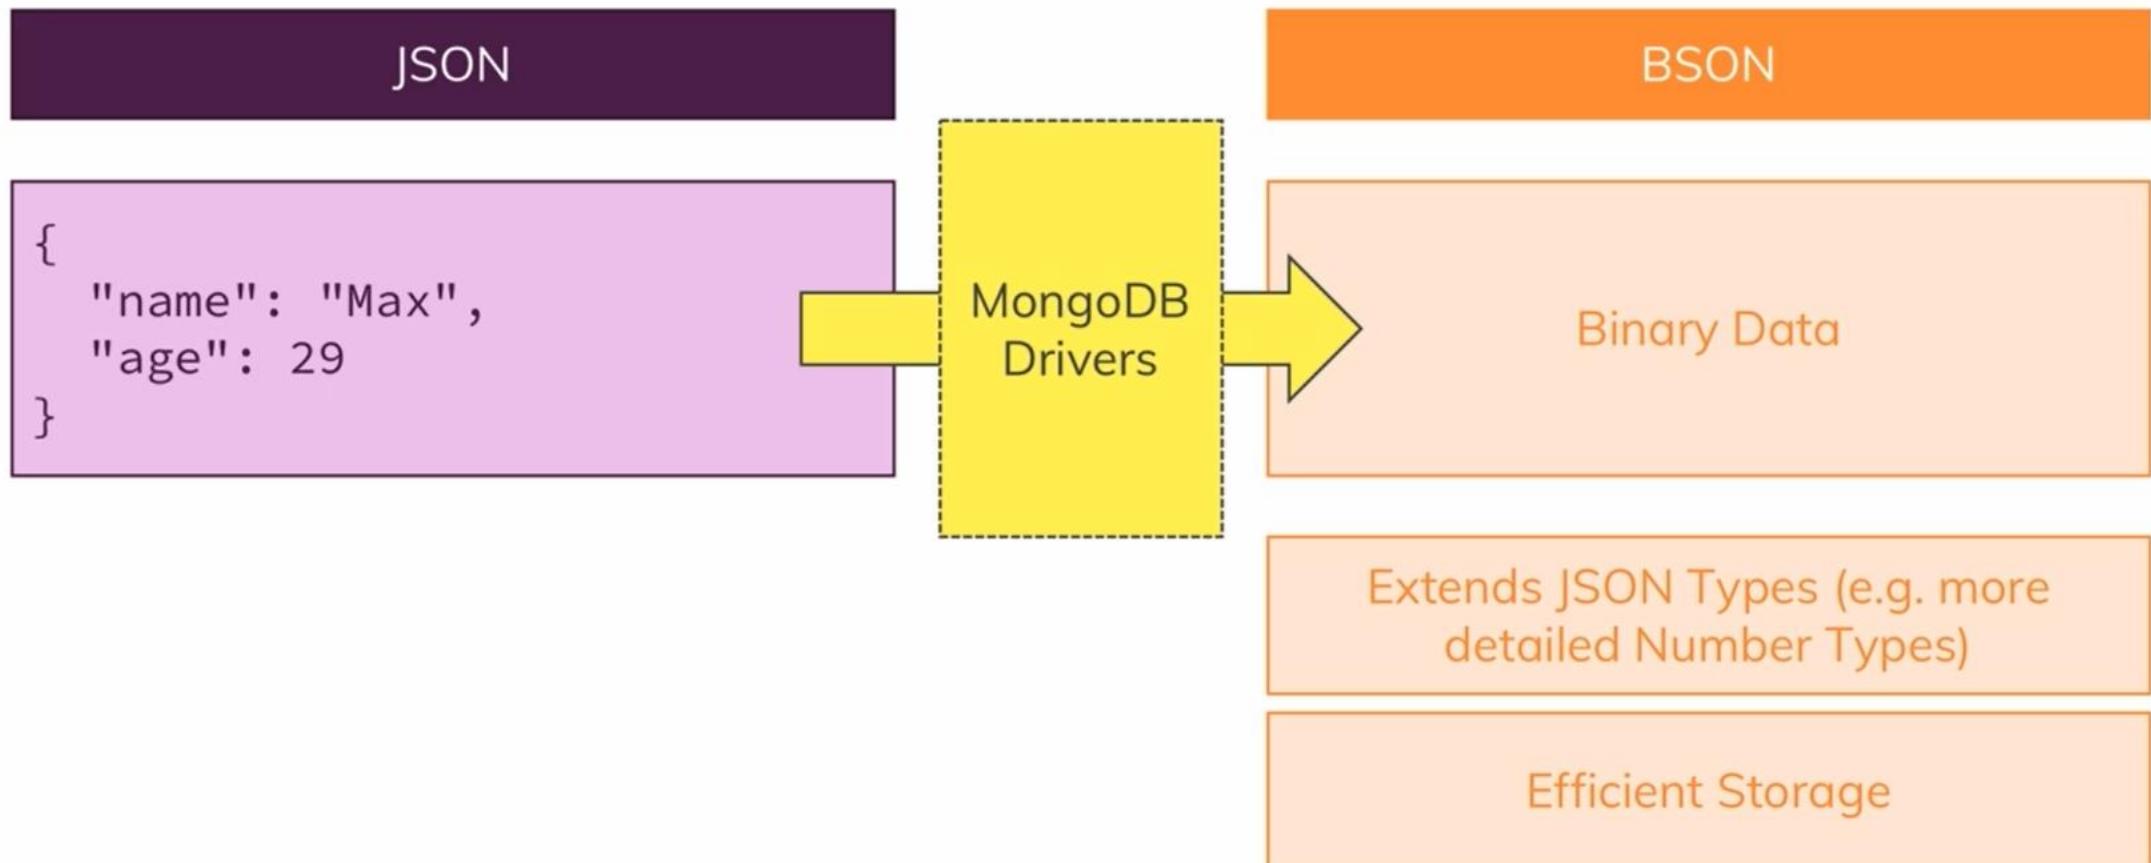
# A Closer Look



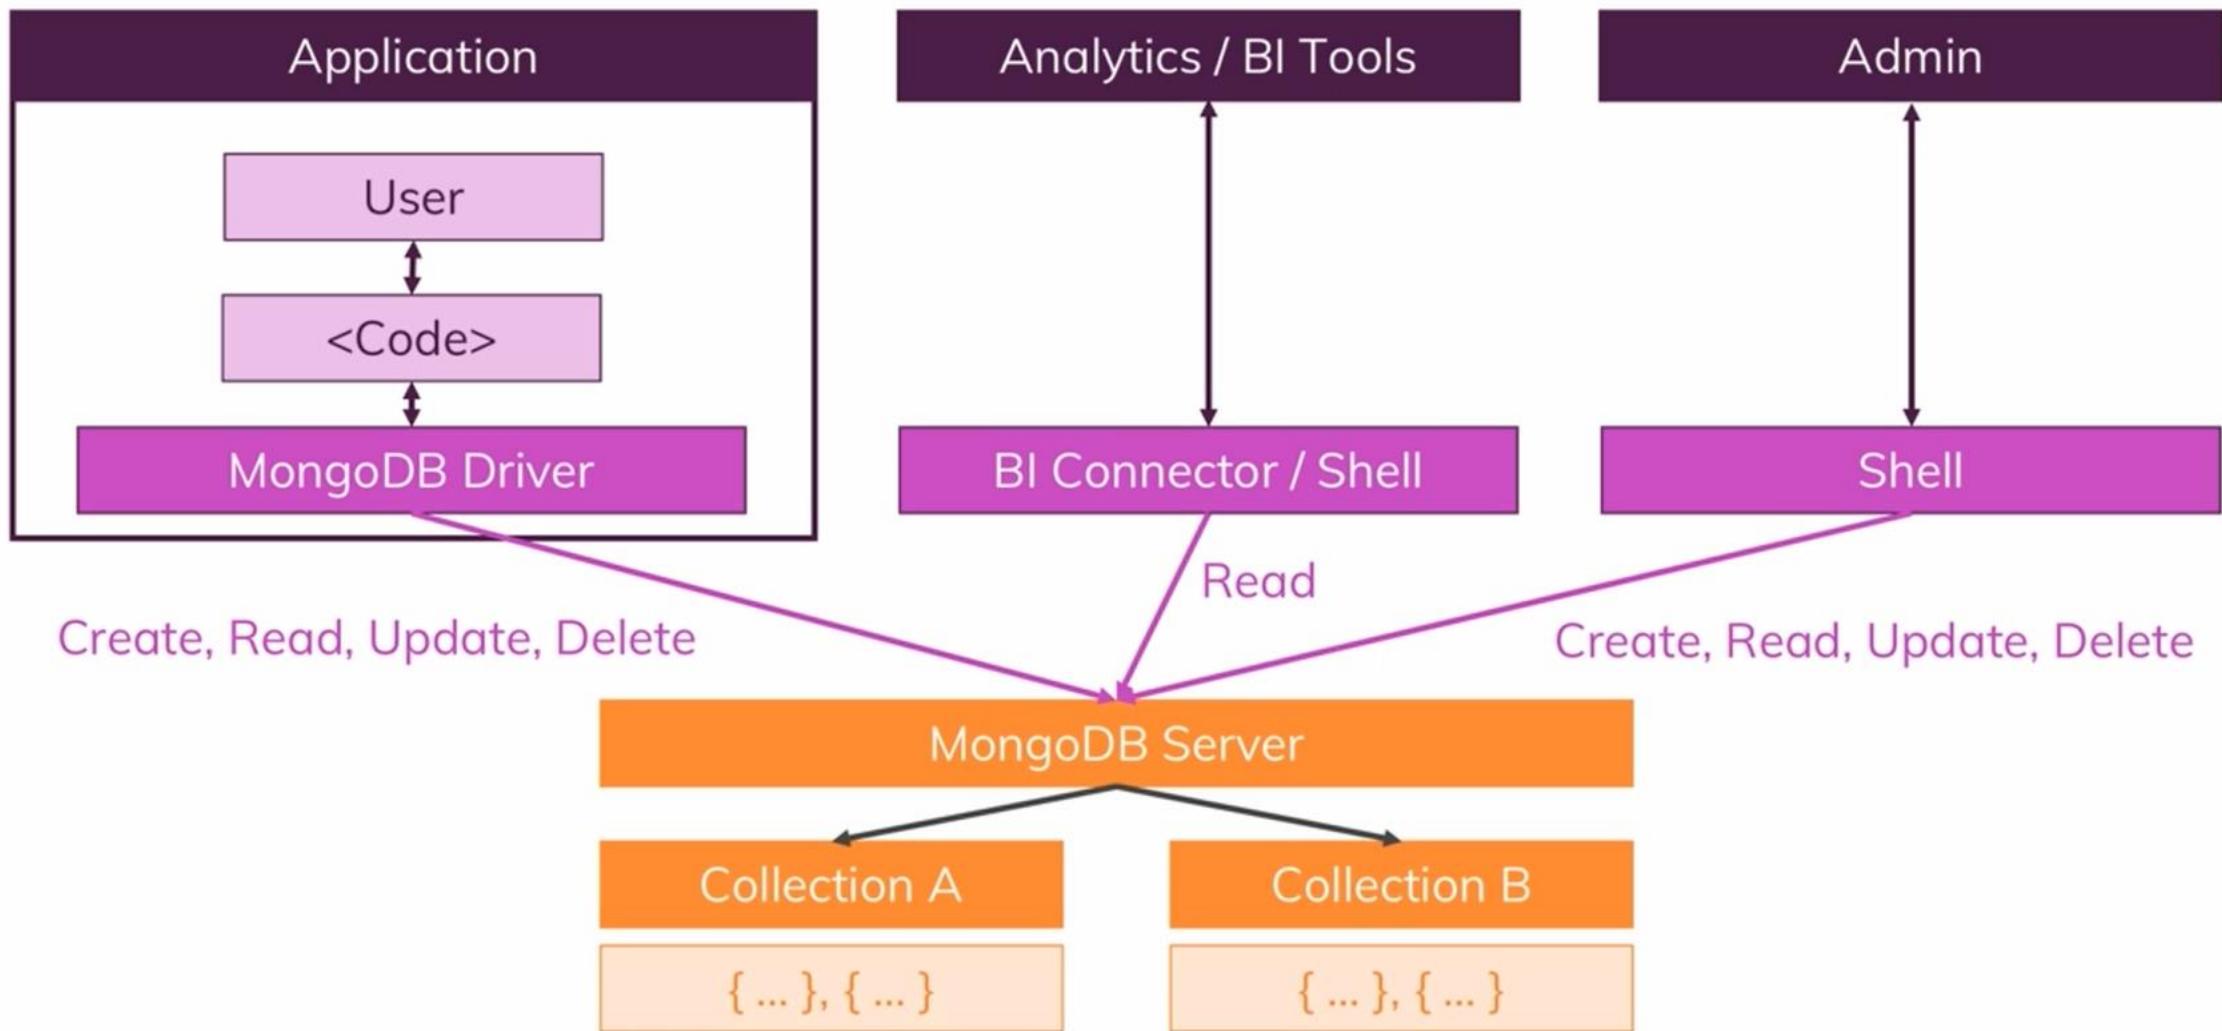
# Databases, Collections, Documents



# JSON vs BSON



# CRUD Operations & MongoDB



# What are CRUD Operations?

CRUD operations refer to the basic Create, Read, Update and Delete operations.

# CRUD - Create, Read, Update, Delete

C	Create or Insert	<code>db.collection.insertOne()</code>
R	Read or Find	<code>db.collection.find()</code>
U	Update	<code>db.collection.updateOne()</code>
D	Delete	<code>db.collection.deleteOne()</code>

# CRUD Operations

## Create

```
insertOne(data, options)
```

```
insertMany(data, options)
```

## Update

```
updateOne(filter, data, options)
```

```
updateMany(filter, data, options)
```

```
replaceOne(filter, data, options)
```

## Read

```
find(filter, options)
```

```
findOne(filter, options)
```

## Delete

```
deleteOne(filter, options)
```

```
deleteMany(filter, options)
```

# Create Operation



`db.collection.insertOne()`

`db.collection.insertMany()`

# Things to Remember for Insert Behavior



All write operations in MongoDB are atomic on the level of a single document

If the collection does not currently exist, insert operations will create the collection

If an inserted document omits the `_id` field, the MongoDB driver automatically generates an `ObjectId` for the `_id` field

# Things to Remember for Find Operation



## **db.collection.find()**

- Query
- Projection
- Read Concern

# Query and Query Operators



**Find method supports many different query operators to filter data**

- Comparison
- Logical
- Element
- Evaluation
- Geospatial
- Array
- Bitwise

# Query Projection



**Specifies the fields to return in the documents that match the query filter**

- 1 or true : include the field
- 0 or false : exclude the field

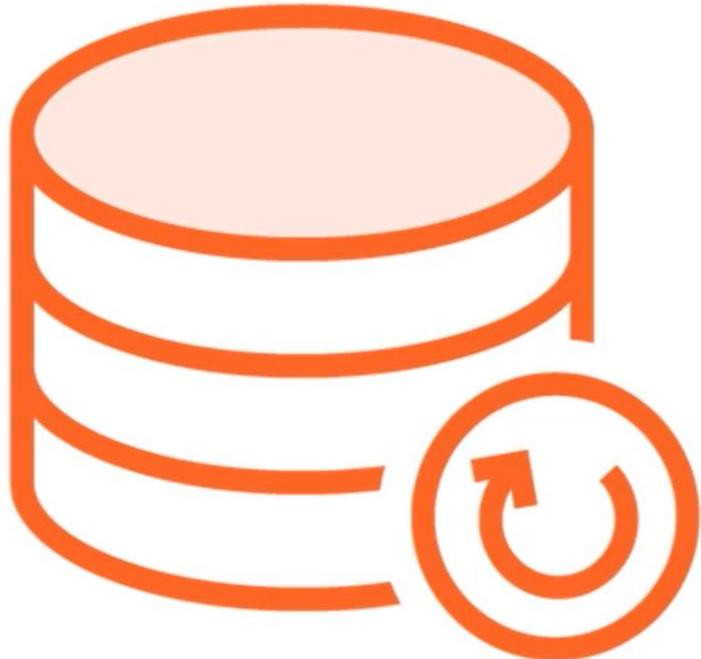
# Read Concern



**Allows to control the consistency and isolation properties of the data read from replica sets and replica set shards**

- Local
- Available
- Majority
- Linearizable
- Snapshot

# Update Operation



**`db.collection.updateOne()`**

**`db.collection.updateMany()`**

**`db.collection.replaceOne()`**

# Write Concern



Level of acknowledgement requested from MongoDB for write operations

w:1 – Ack from primary

w:0 – No ack

w:(n) – Primary + (n-1) secondary

w: majority

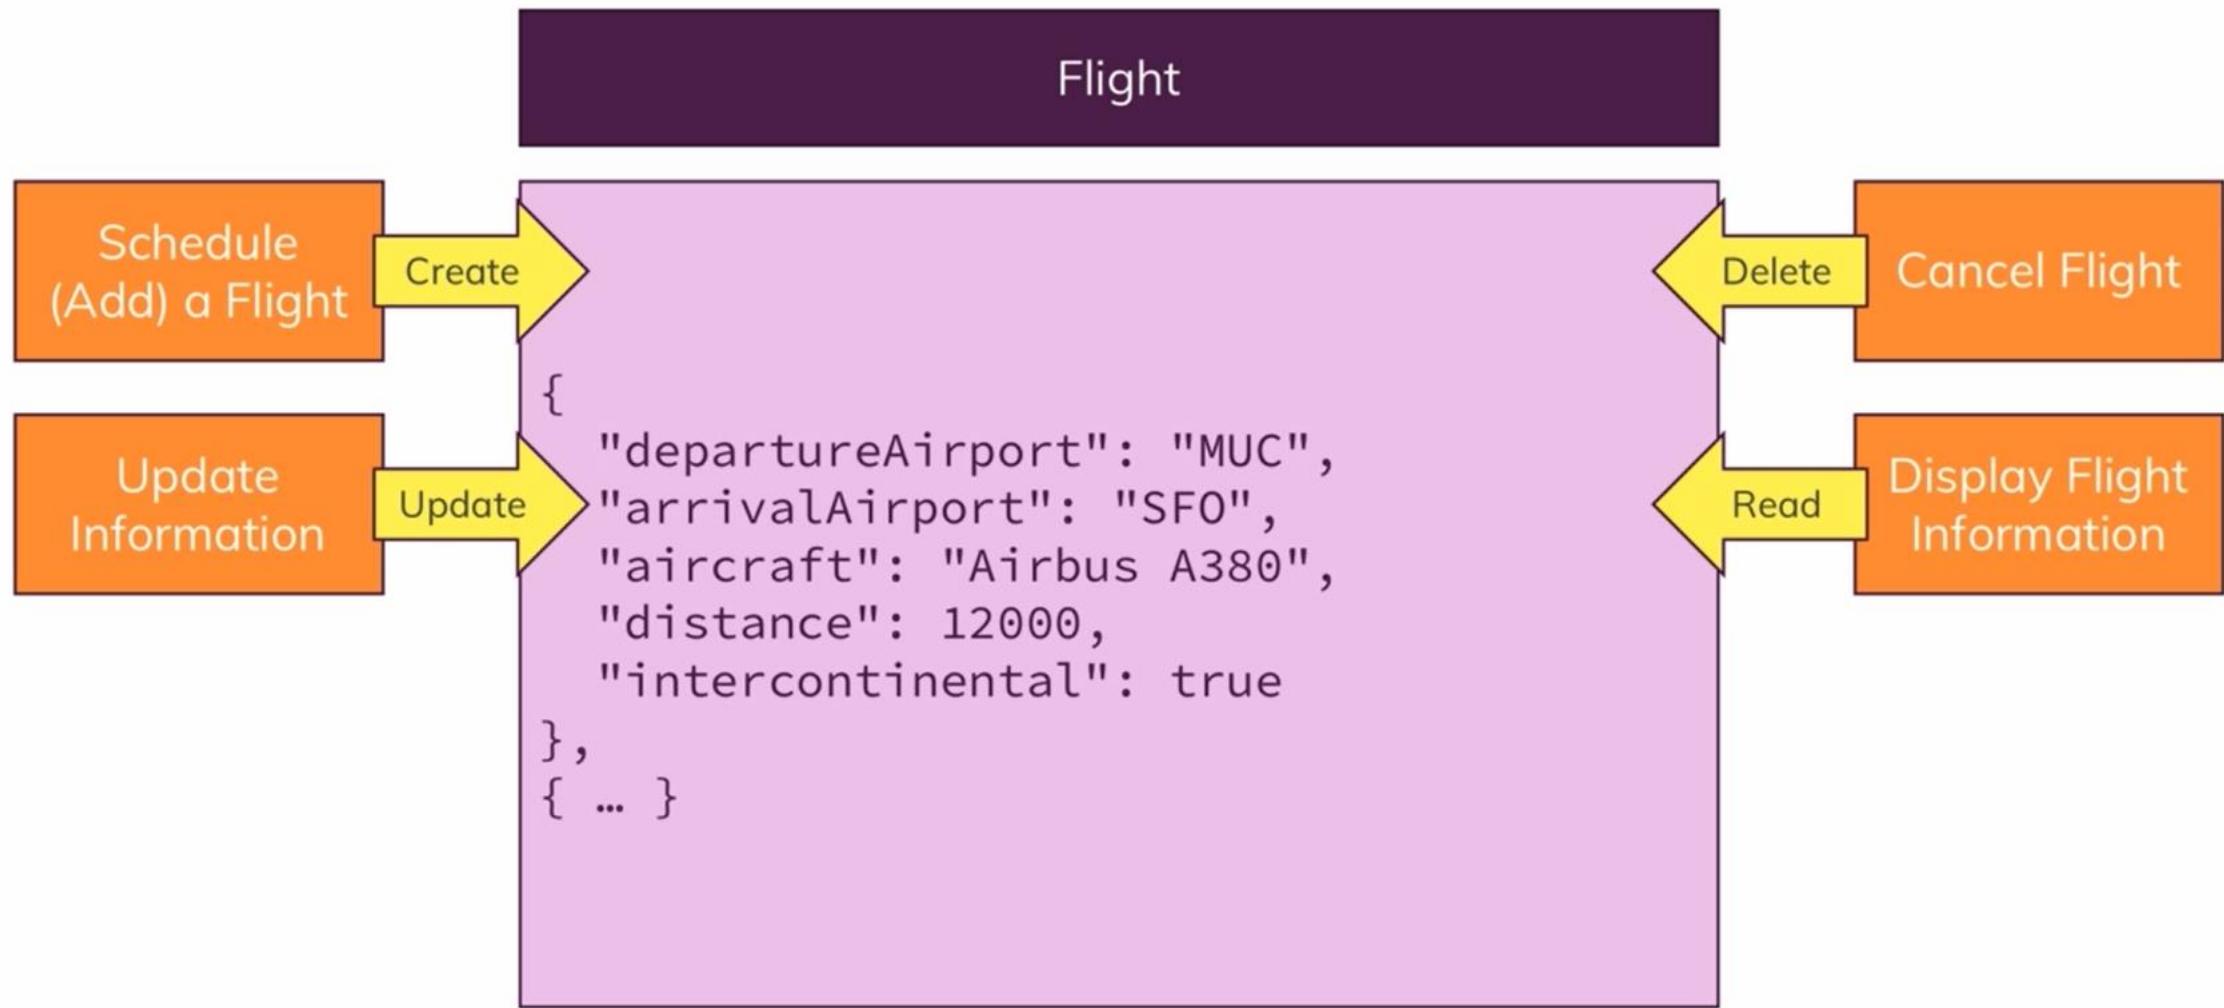
wtimeout: Time limit to prevent write operations from blocking indefinitely

# Things to Remember for Update Behavior

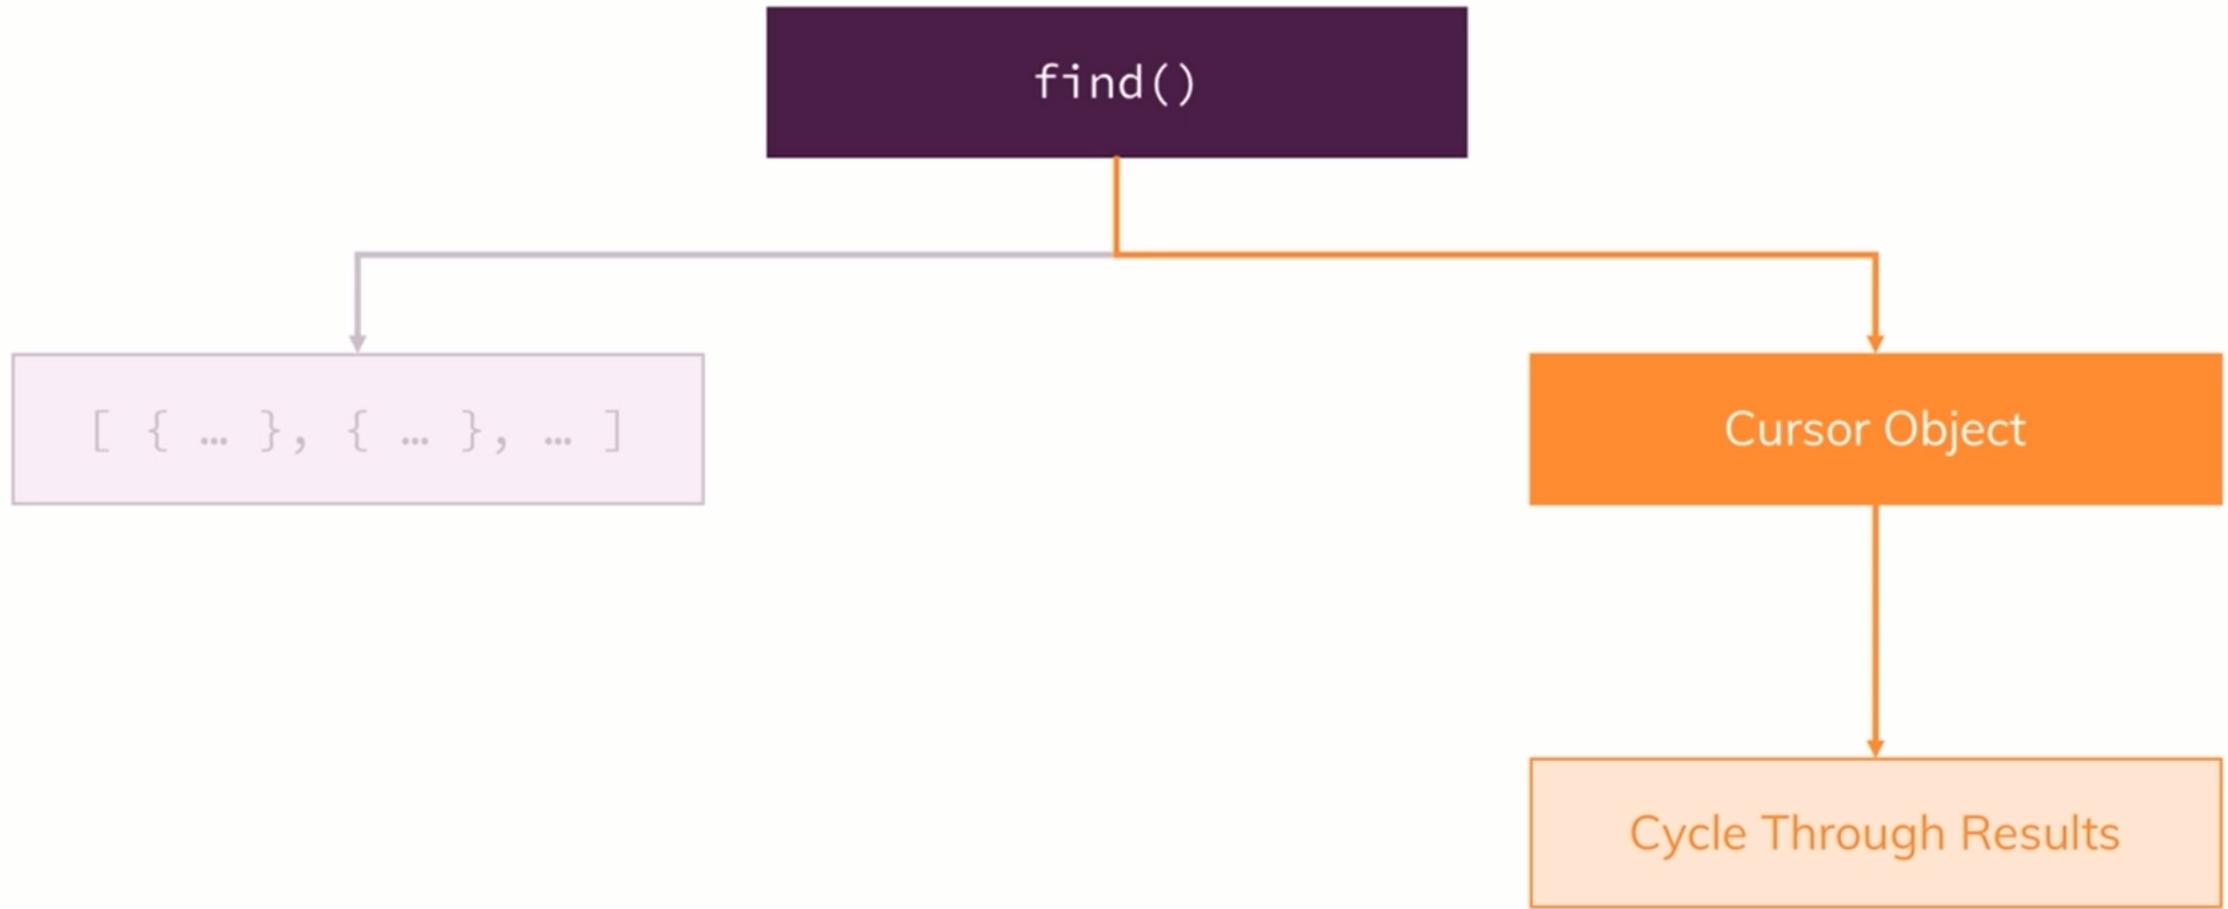


- Atomic on the level of a single document**
- \_id field cannot be replaced with different value**
- \$set creates field if not already existing**
- upsert : true**
  - Update on match of filter
  - Insert no match of filter

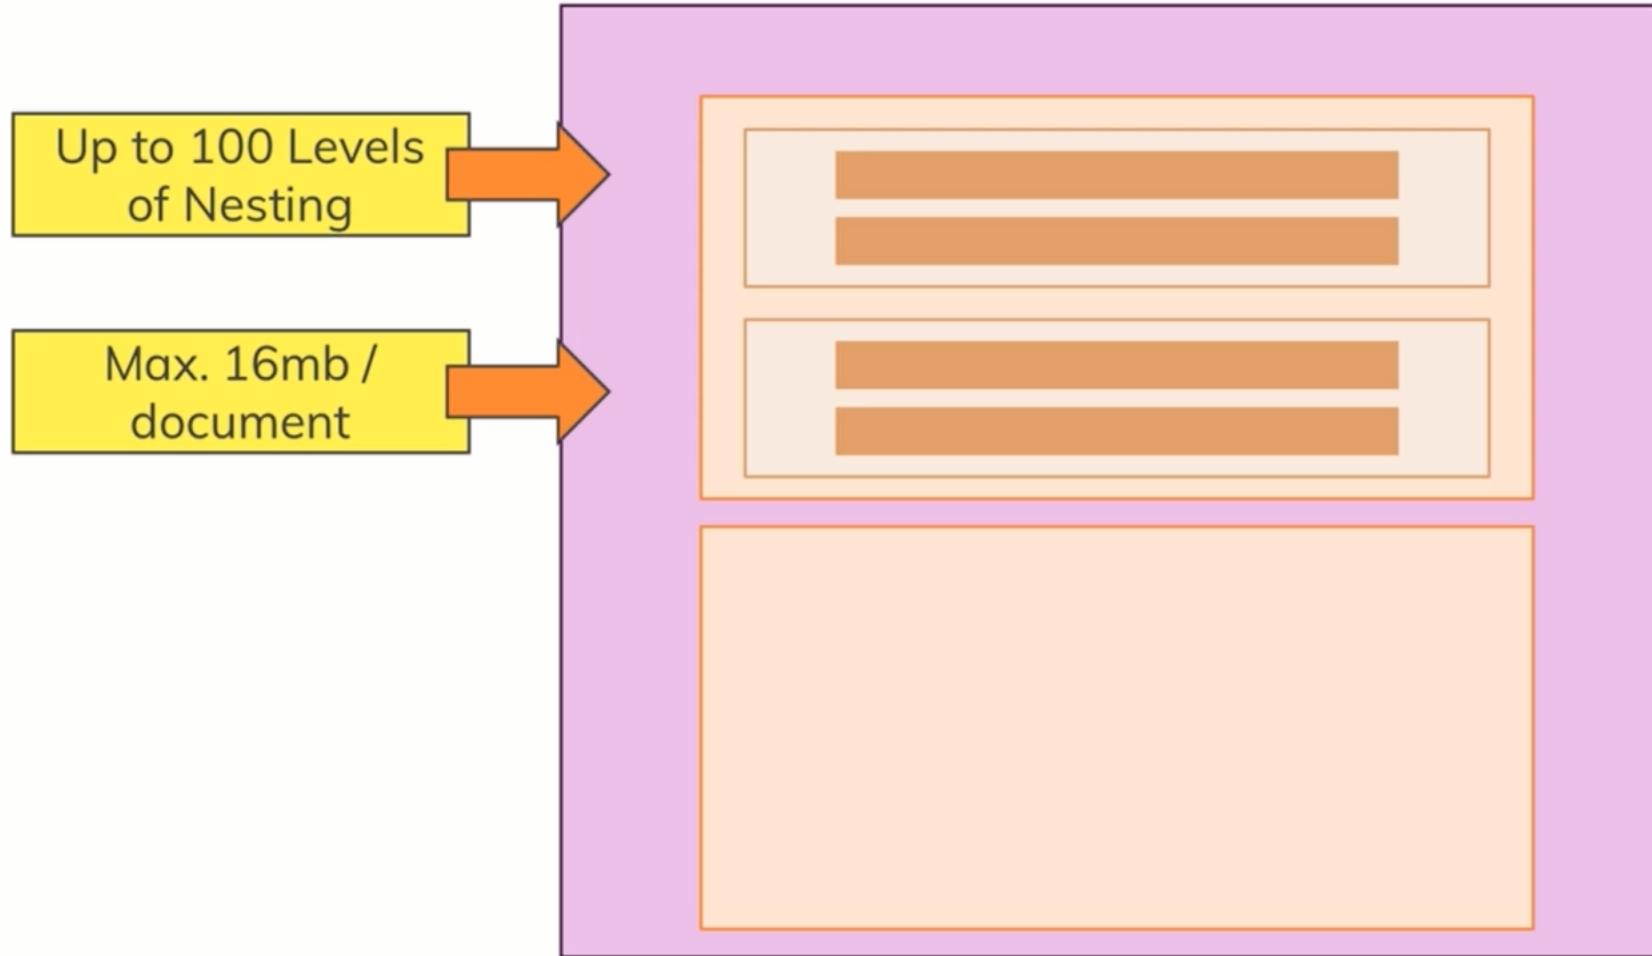
# Example #1: Flight Data



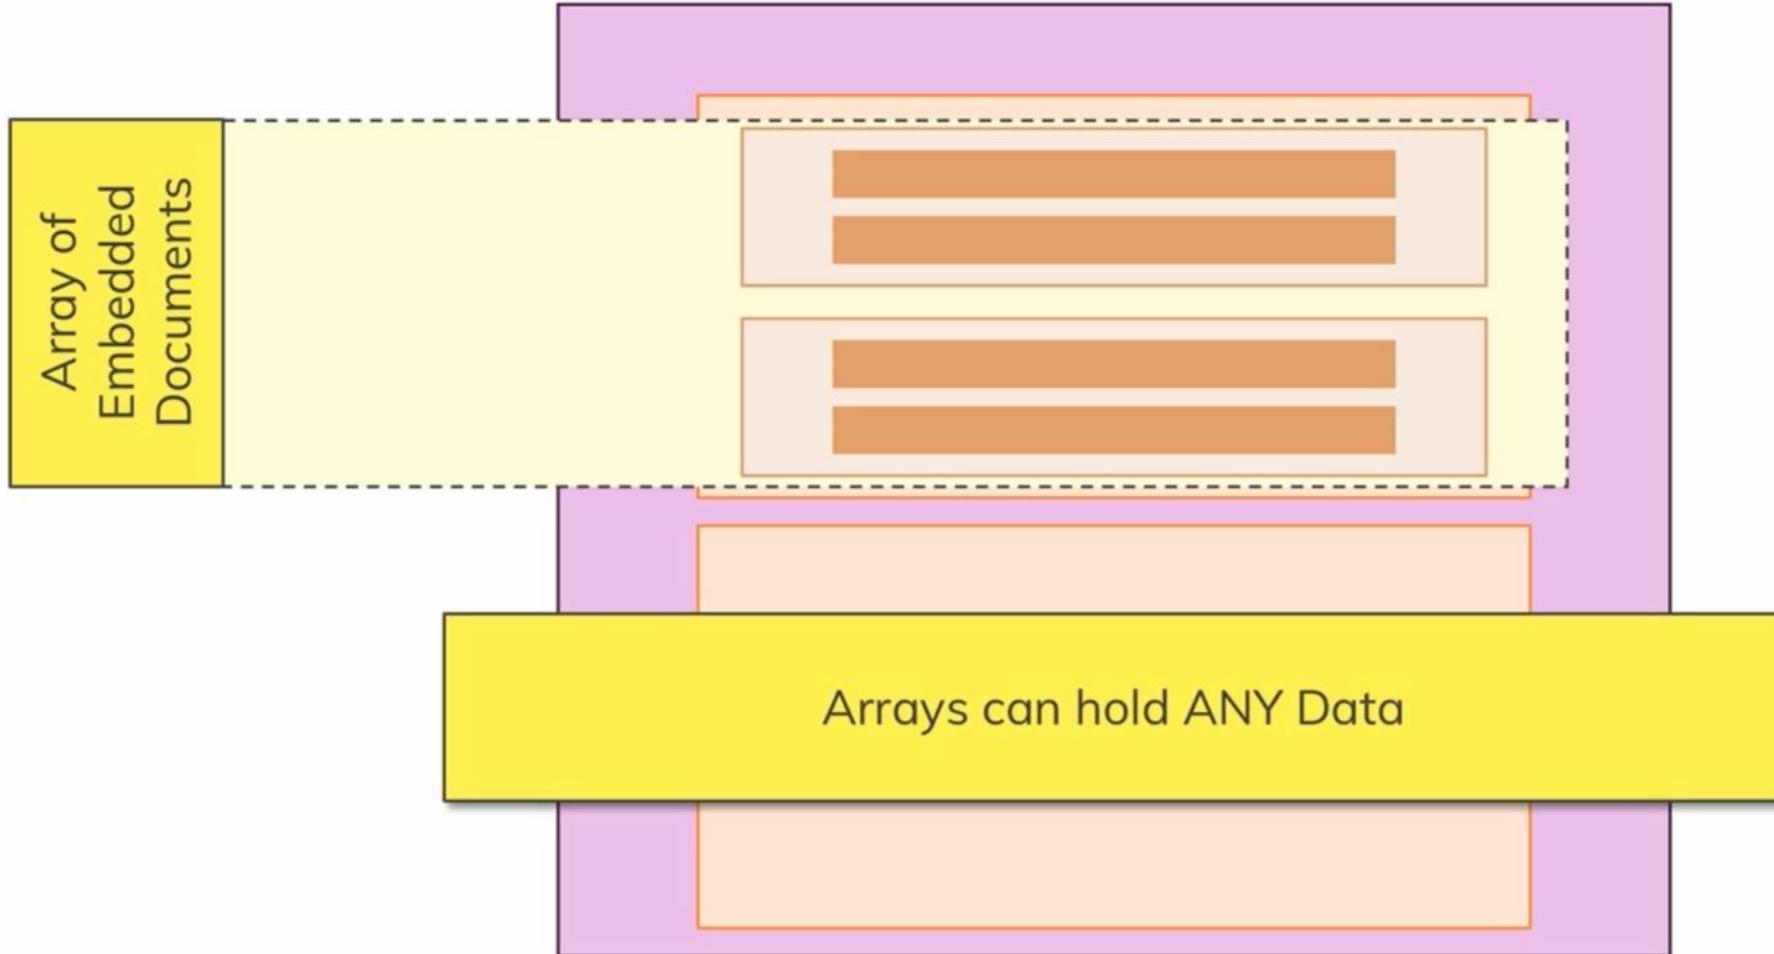
# Cursors



# Embedded Documents



# Arrays



## Databases, Collections, Documents

- A Database holds multiple Collections where each Collection can then hold multiple Documents
- Databases and Collections are created “lazily” (i.e. when a Document is inserted)
- A Document can’t directly be inserted into a Database, you need to use a Collection!

## CRUD Operations

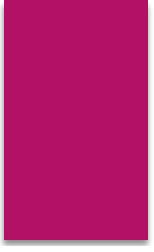
- CRUD = Create, Read, Update, Delete
- MongoDB offers multiple CRUD operations for single-document and bulk actions (e.g. `insertOne()`, `insertMany()`, ...)
- Some methods require an argument (e.g. `insertOne()`), others don’t (e.g. `find()`)
- `find()` returns a cursor, NOT a list of documents!
- Use filters to find specific documents

## Document Structure

- Each document needs a unique ID (and gets one by default)
- You may have embedded documents and array fields

## Retrieving Data

- Use filters and operators (e.g. `$gt`) to limit the number of documents you retrieve



Understanding Document Schemas &  
Data Types

Modelling Relations

Schema Validation

# Schema-less Or Not?

Isn't MongoDB all about having **NO** data Schemas?



MongoDB enforces no schemas! Documents don't have to use the same schema inside of one collection



But that does not mean that you can't use some kind of schema!

# To Schema Or Not To Schema

Chaos!



SQL World!

Products

```
{  
  "title": "Book",  
  "price": 12.99  
}
```

Very Different!

```
{  
  "name": "Bottle",  
  "available": true  
}
```

Products

```
{  
  "title": "Book",  
  "price": 12.99  
}
```

Extra Data

```
{  
  "title": "Bottle",  
  "price": 5.99  
  "available": true  
}
```

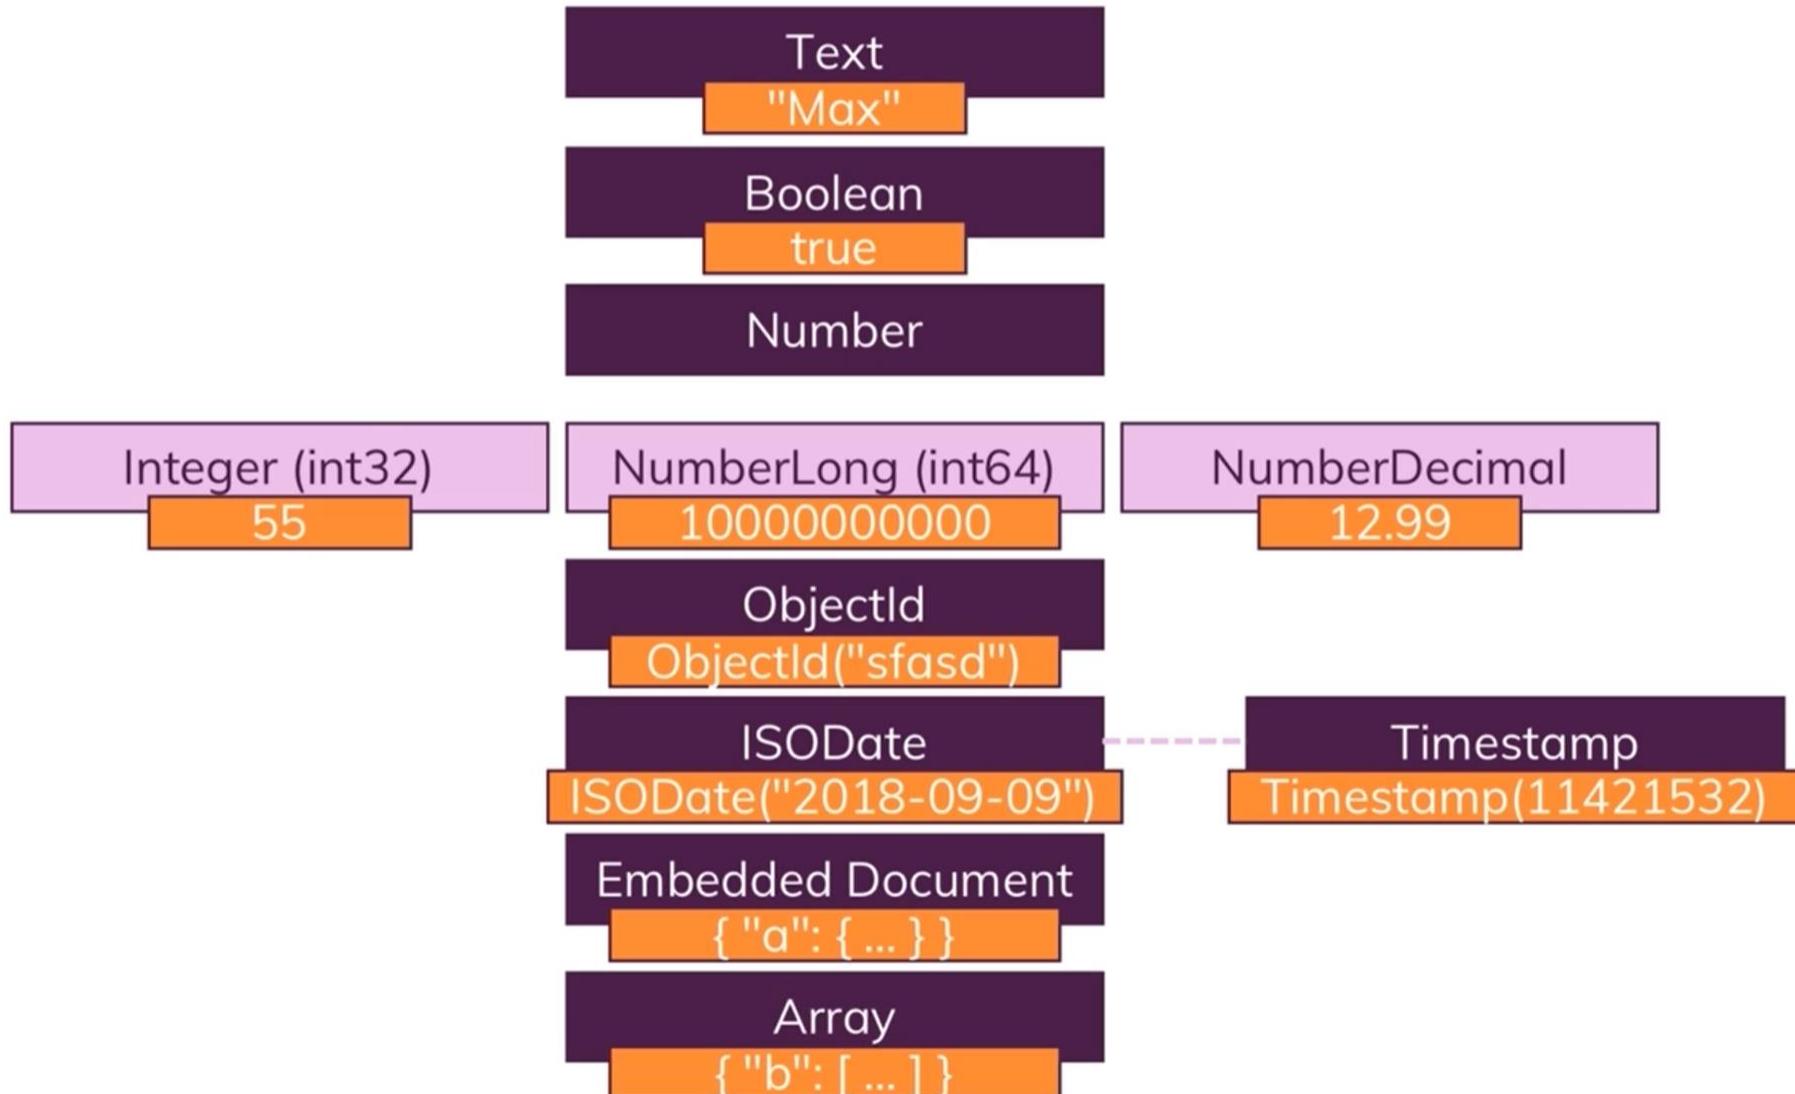
Products

```
{  
  "title": "Book",  
  "price": 12.99  
}
```

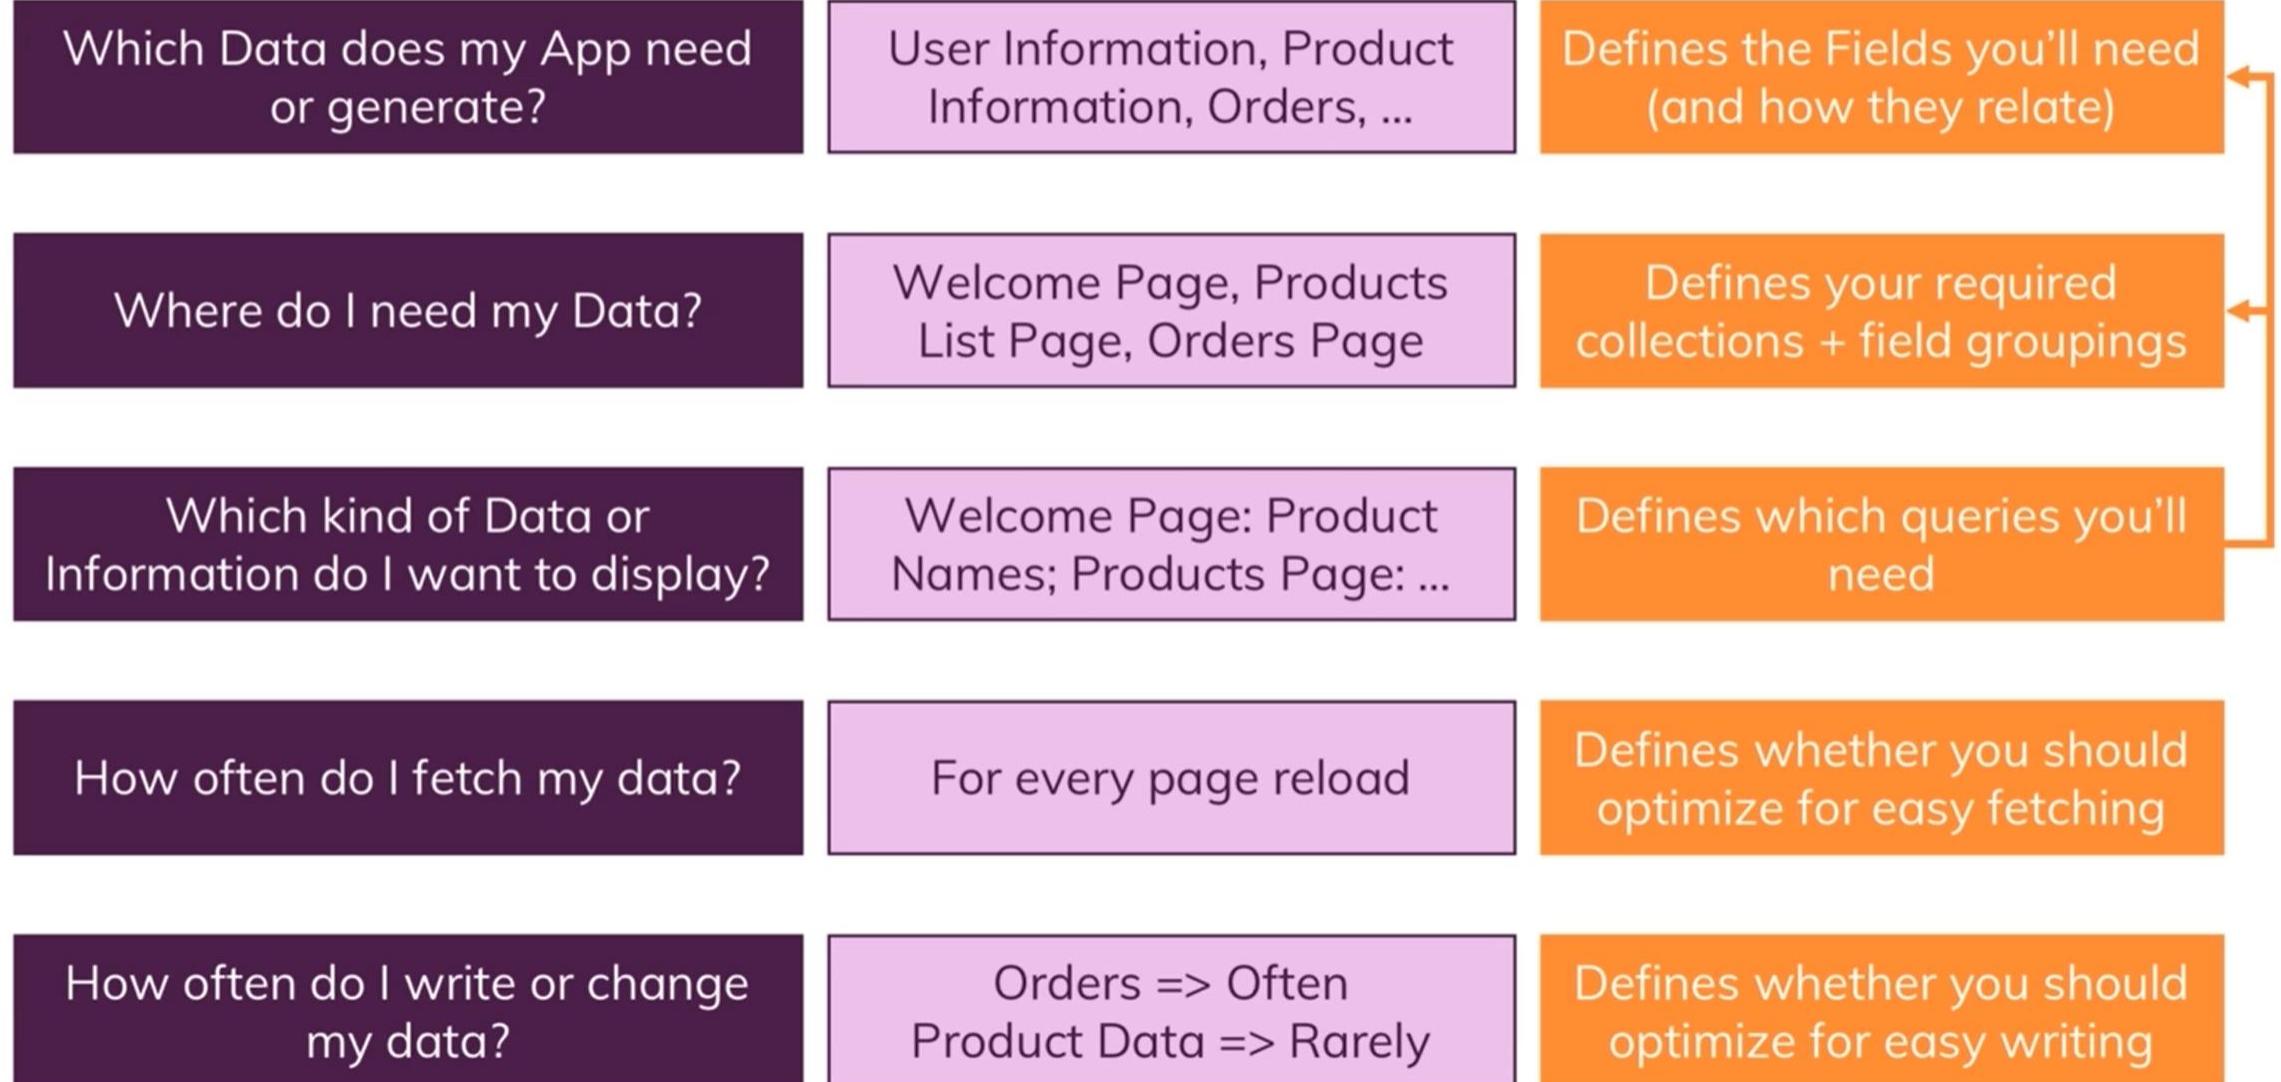
Full Equality

```
{  
  "title": "Bottle",  
  "price": 5.99  
}
```

# Data Types



# Data Schemas & Data Modelling



# Relations - Options

## Nested / Embedded Documents

```
Customers  
{  
  userName: 'max',  
  age: 29,  
  address: {  
    street: 'Second Street',  
    city: 'New York'  
  }  
}
```

## References

```
{  
  userName: 'max',  
  favBooks: ['id1', 'id2']  
}  
  
Lots of data duplication!
```

```
Customers  
{  
  userName: 'max',  
  favBooks: ['id1', 'id2']  
}
```

```
Books  
{  
  _id: 'id1',  
  name: 'Lord of the Rings 1'  
}
```

## Example #1 – Patient <-> Disease Summary



"One patient has one disease summary, a disease summary belongs to one patient"



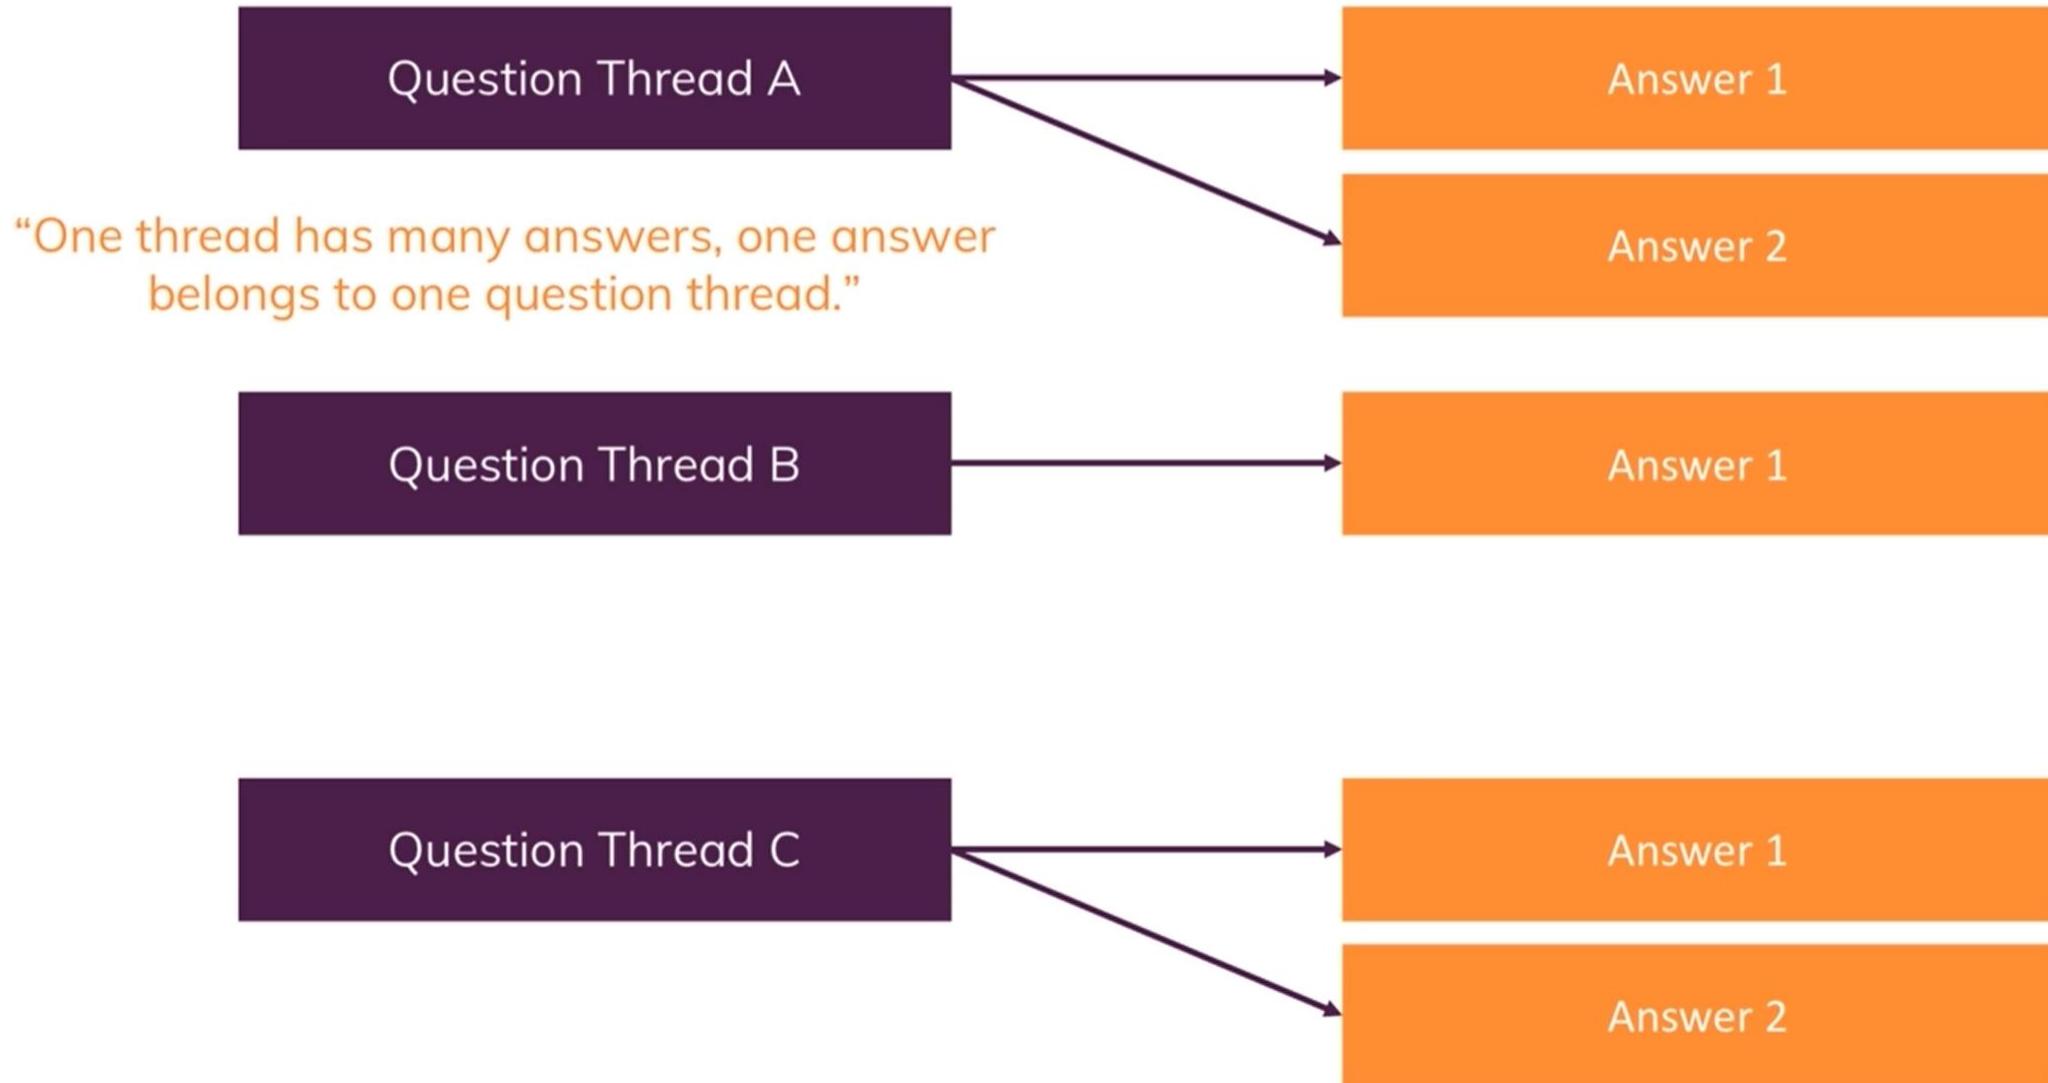
## Example #2 – Person <-> Car



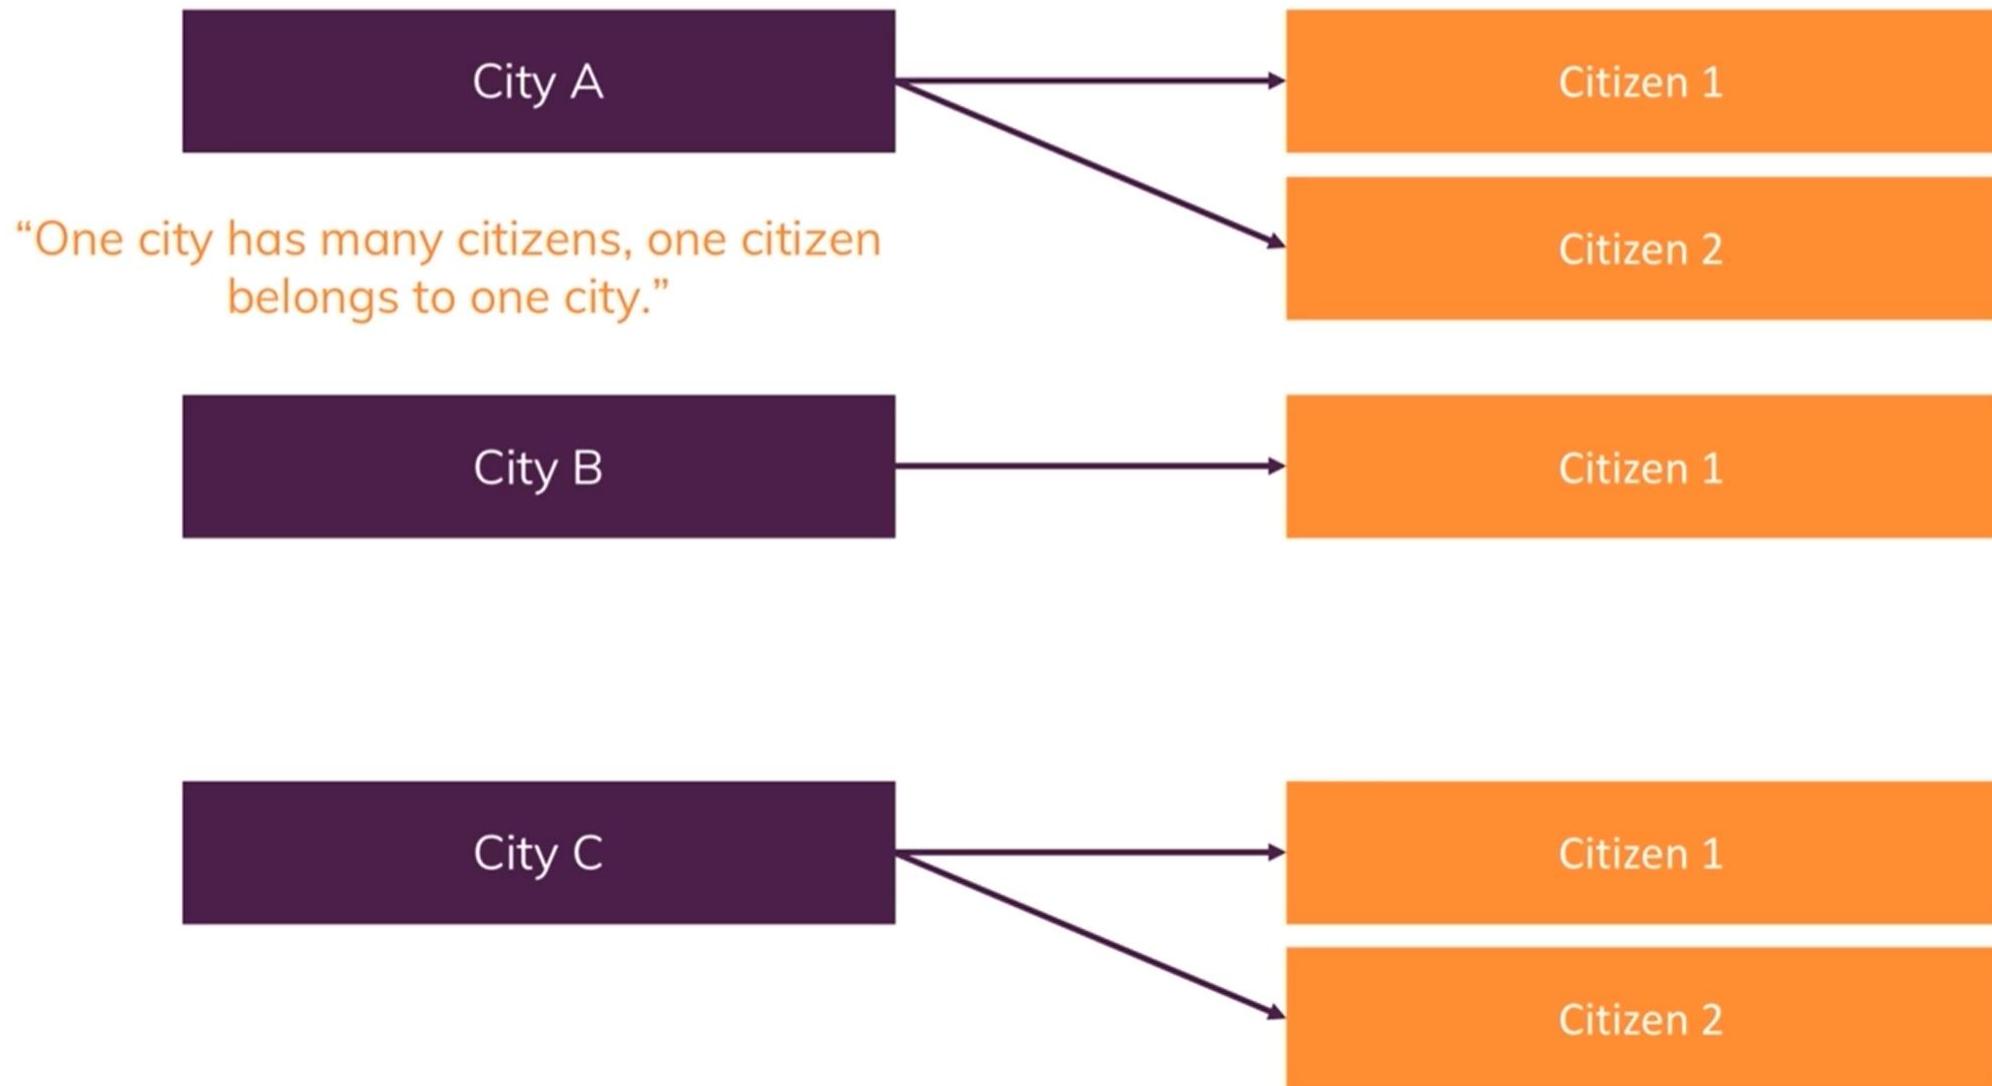
"One person has one car, a car belongs to one person"



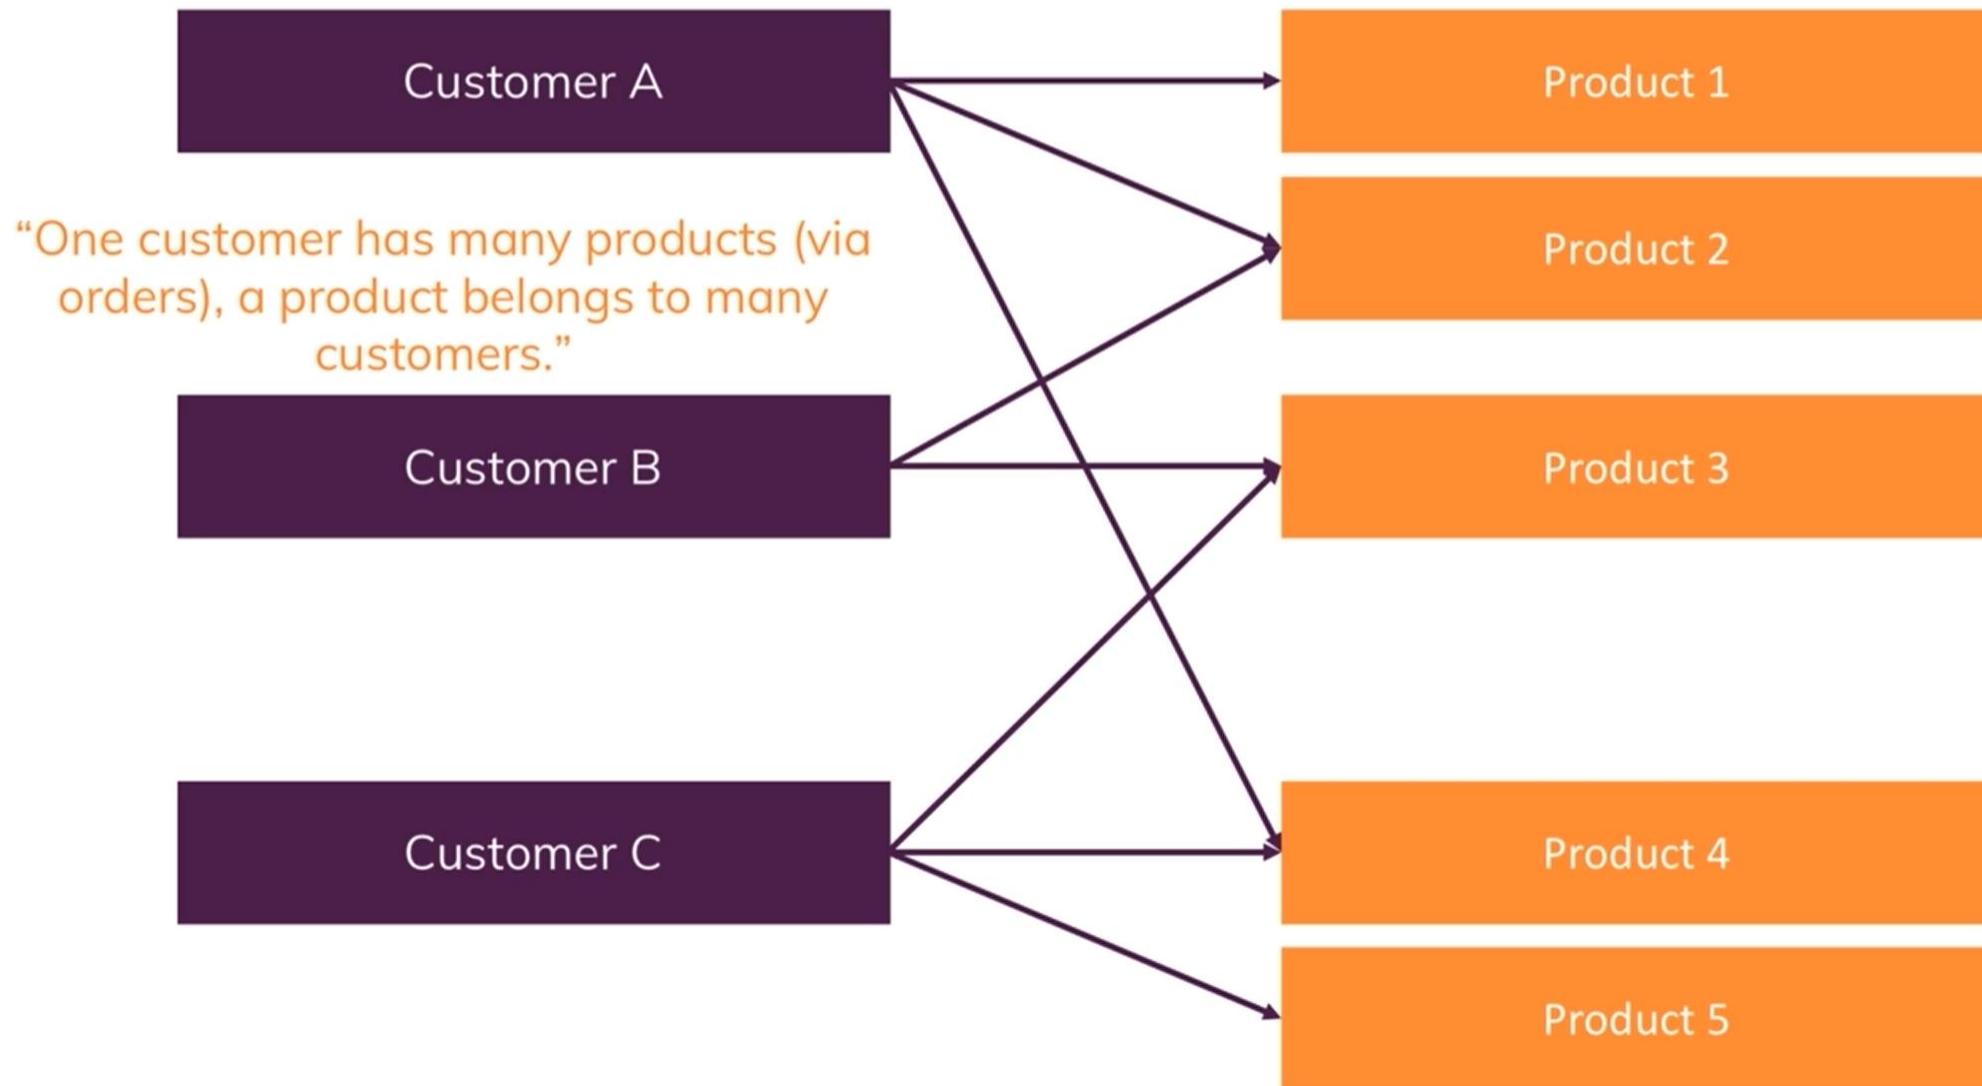
## Example #3 – Thread <-> Answers



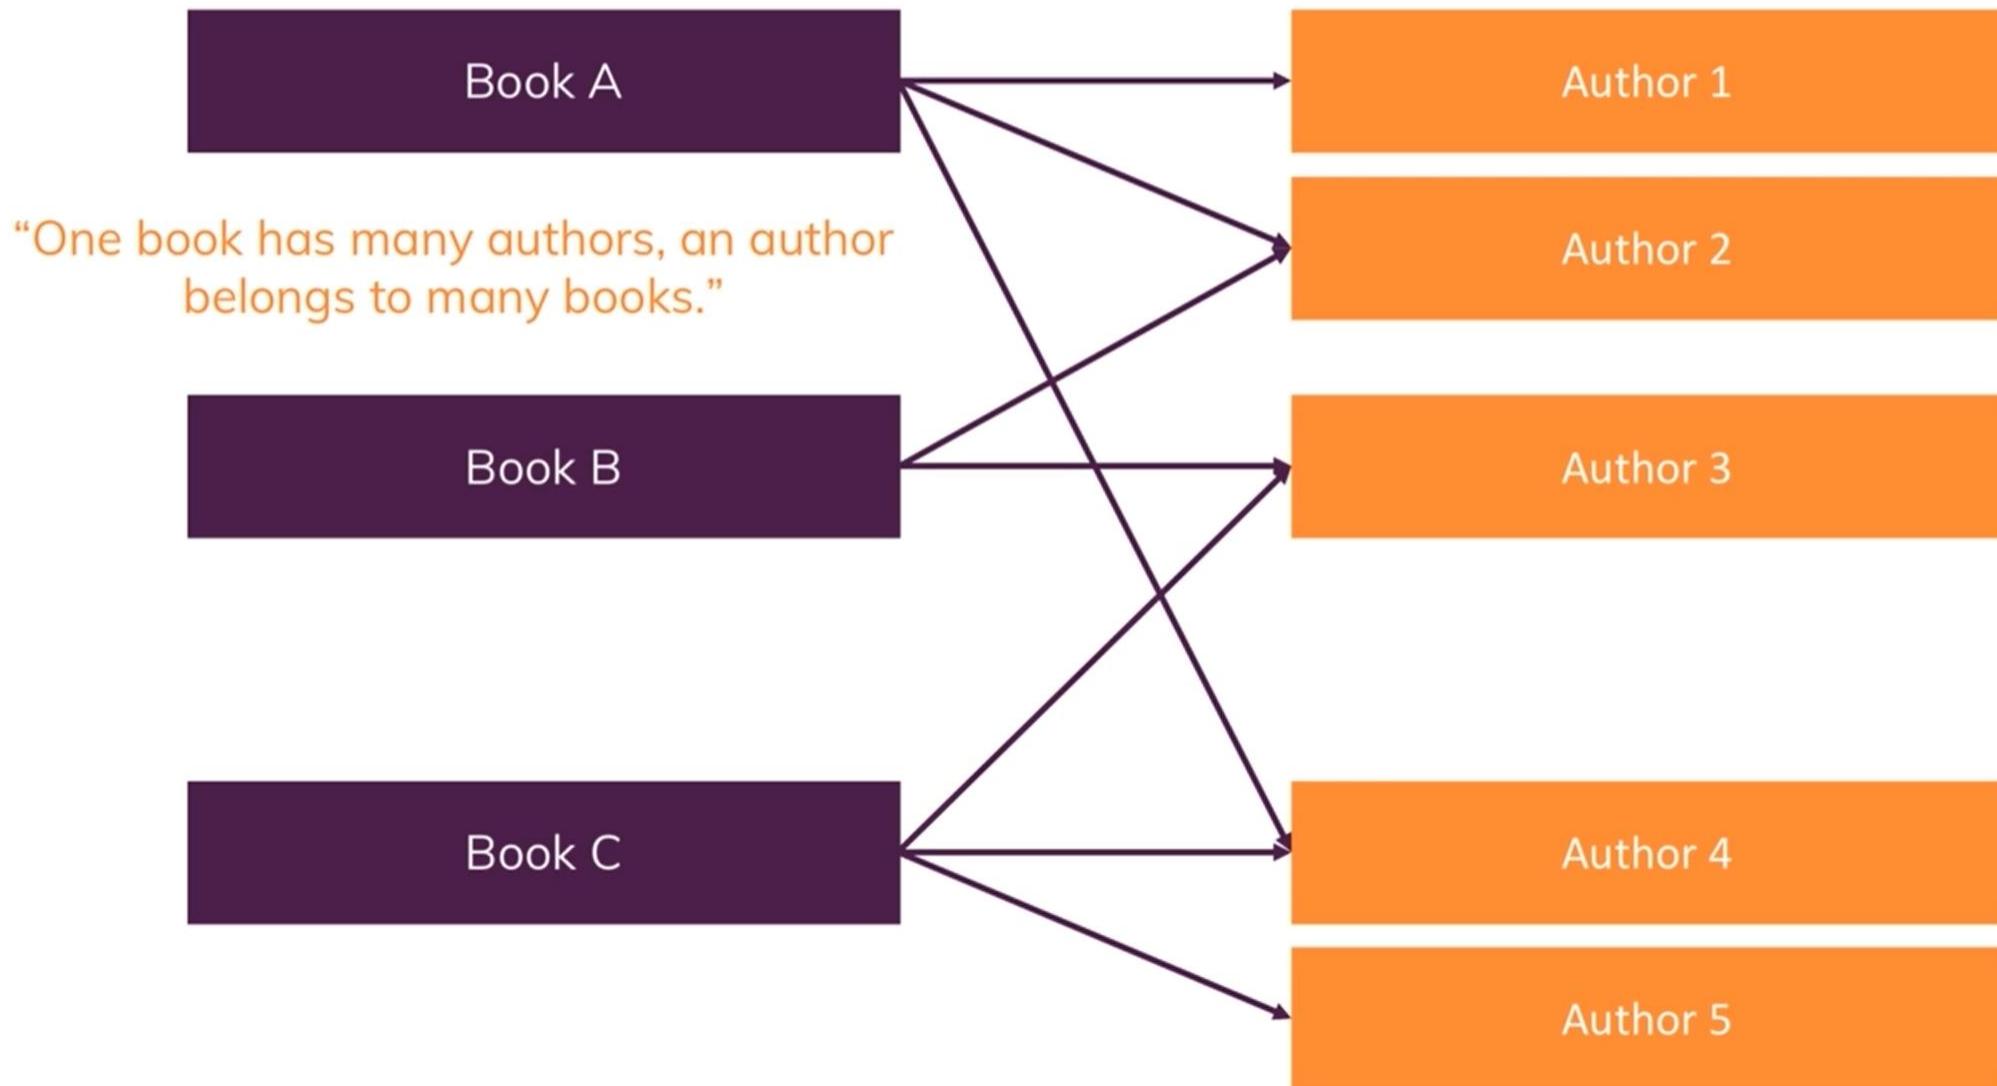
## Example #4 – City <-> Citizens



## Example #5 – Customers <-> Products (Orders)



## Example #6 – Books <-> Authors



# Relations - Options

## Nested / Embedded Documents

Group data together logically

Great for data that belongs together and is not really overlapping with other data

Avoid super-deep nesting (100+ levels) or extremely long arrays (16mb size limit per document)

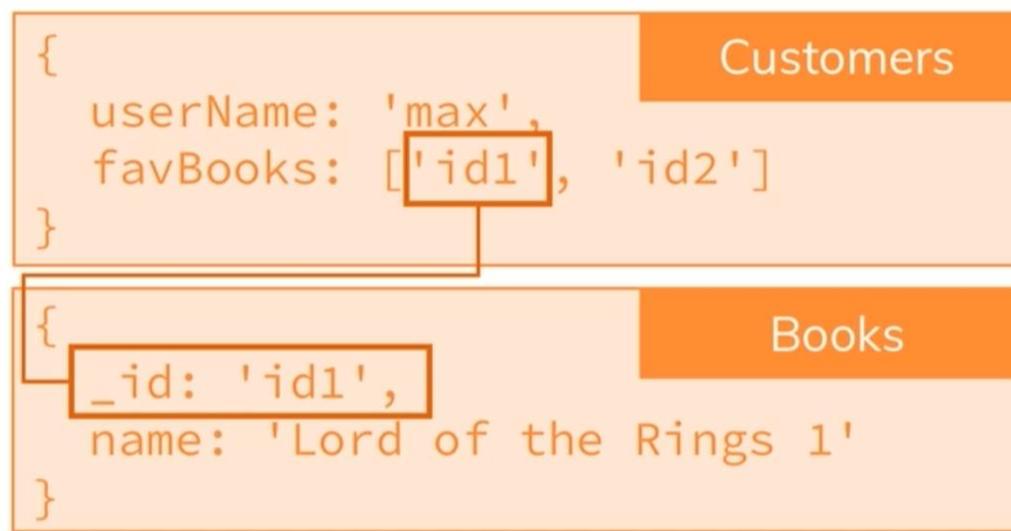
## References

Split data across collections

Great for related but shared data as well as for data which is used in relations and standalone

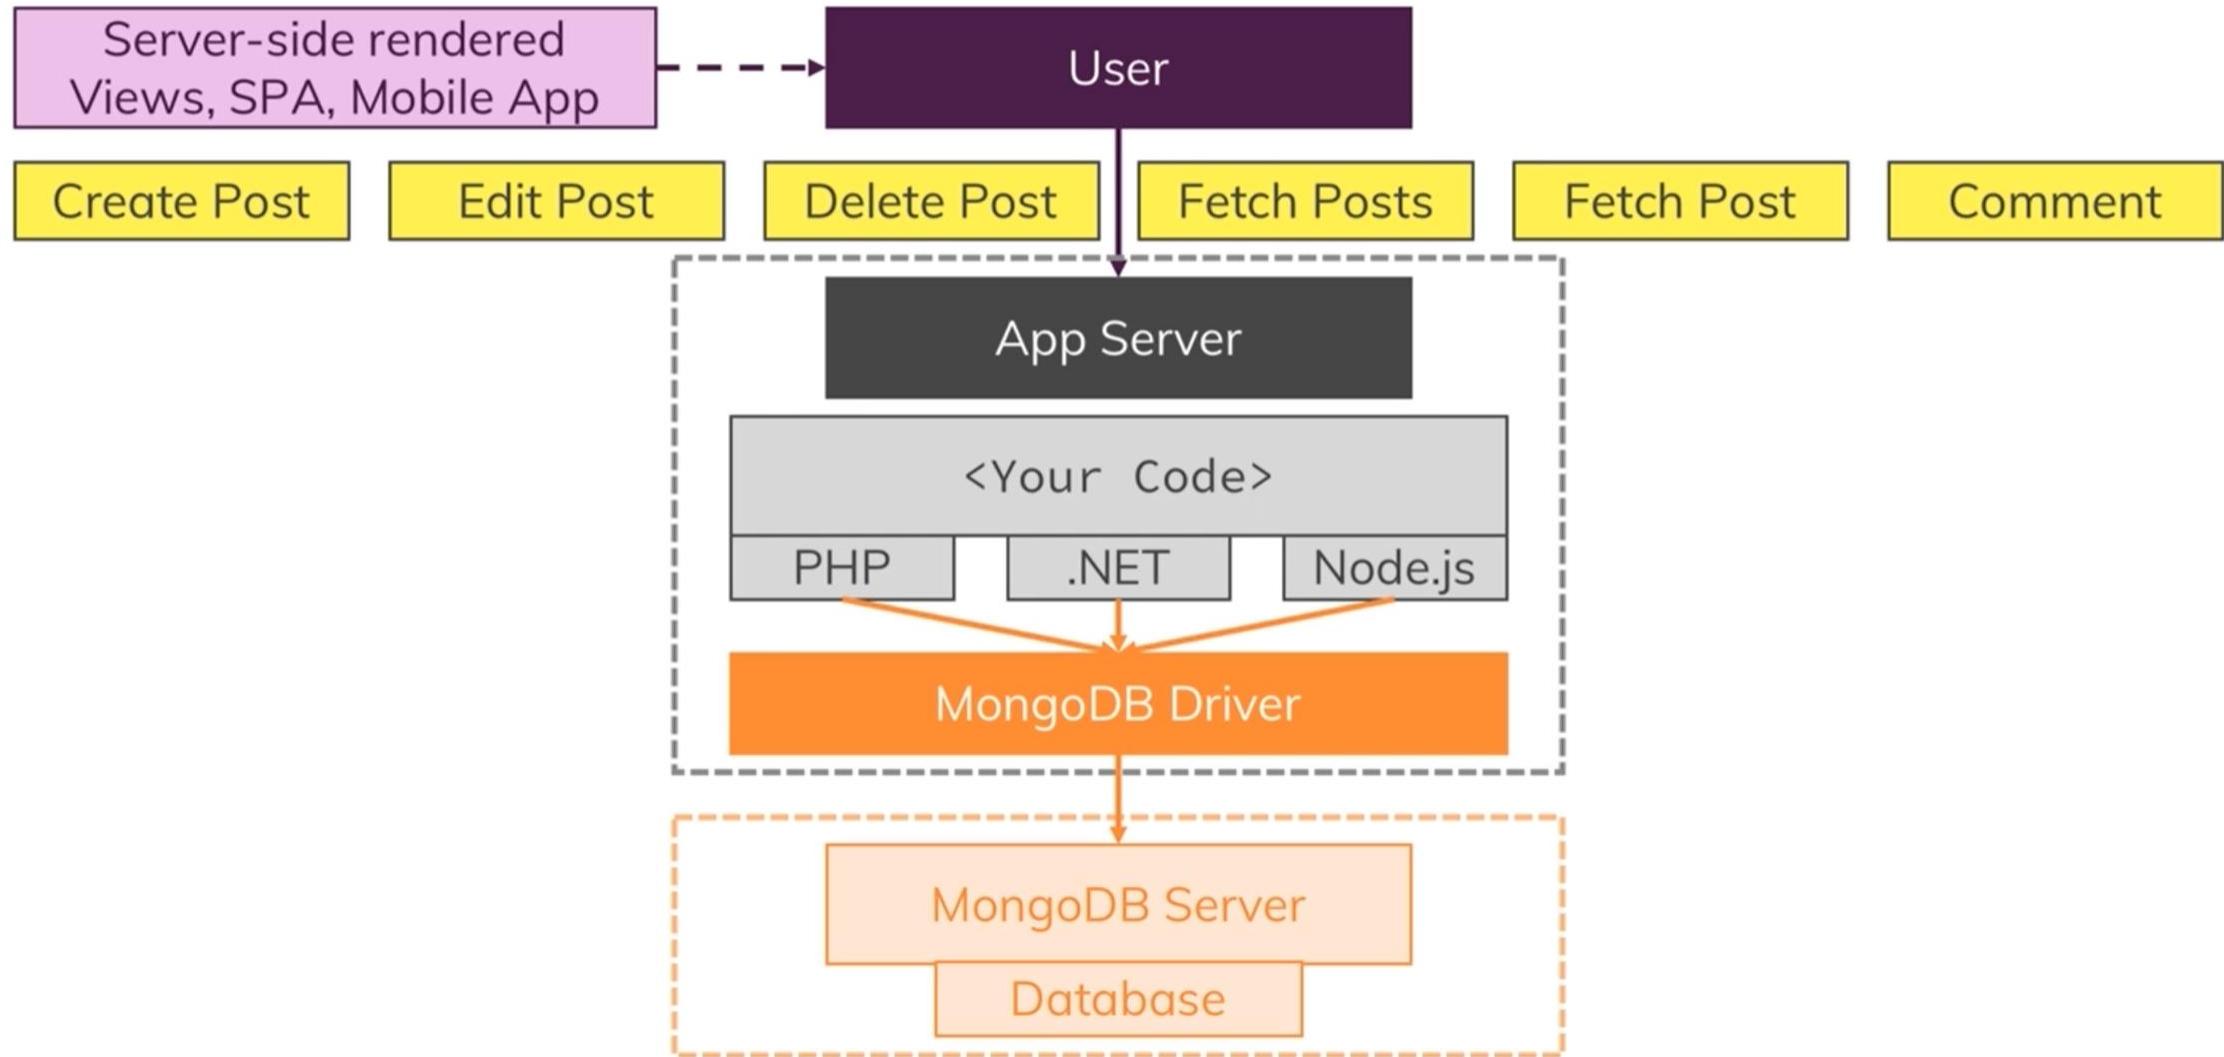
Allows you to overcome nesting and size limits (by creating new documents)

# Joining with \$lookup

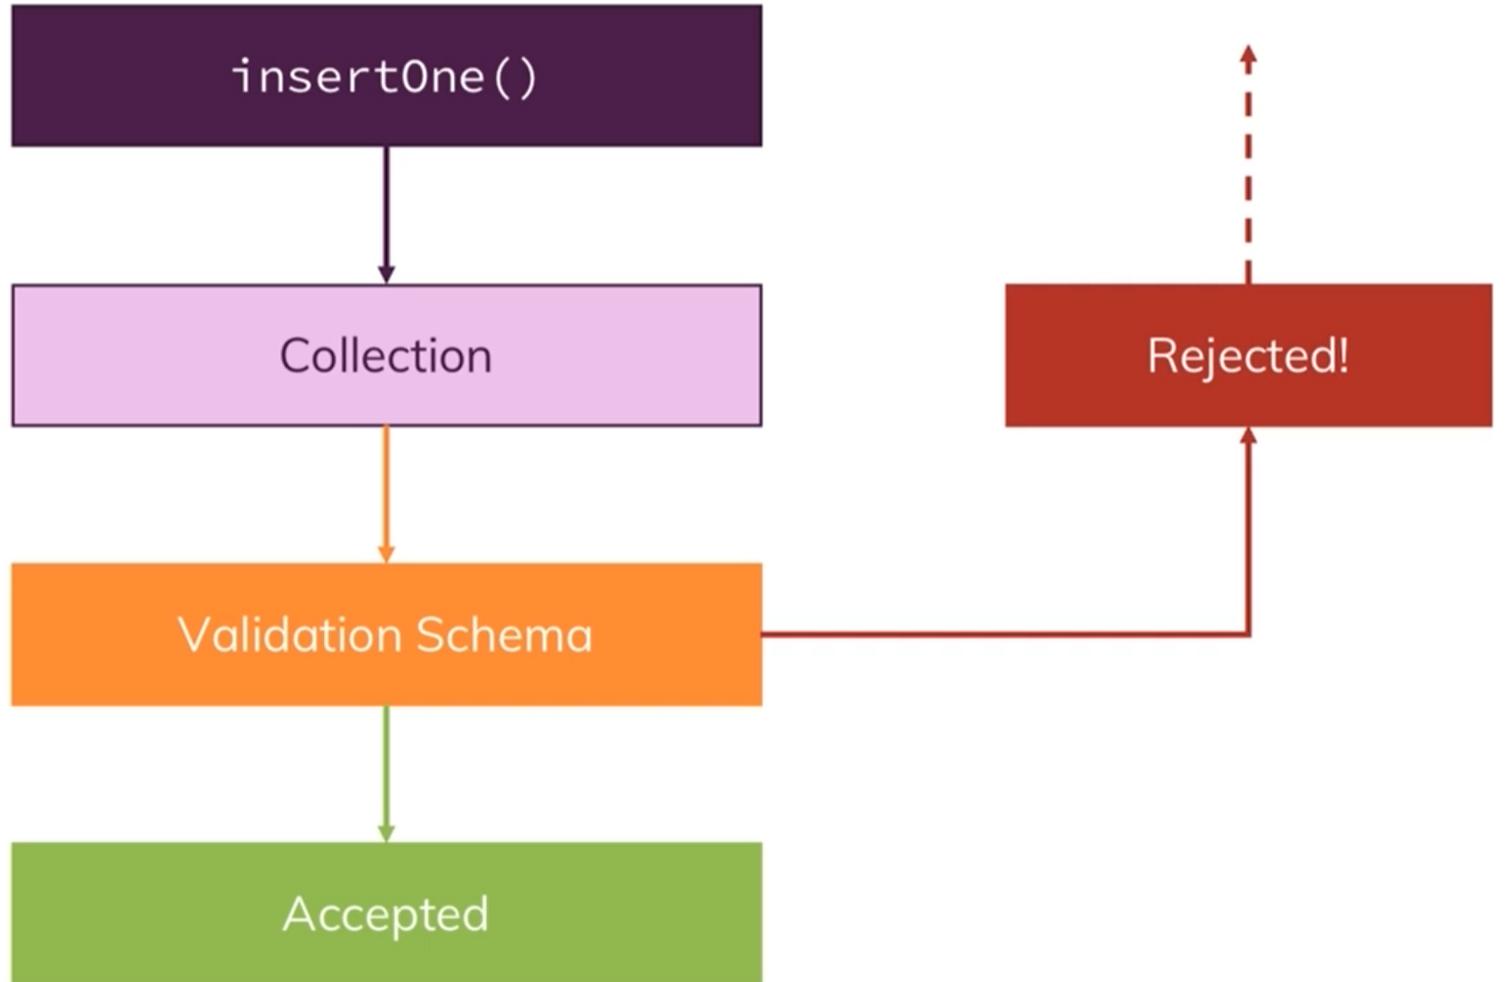


```
customers.aggregate([
  { $lookup: {
      from: "books",
      localField: "favBooks",
      foreignField: "_id"
      as: "favBookData"
    }
  }
])
```

# Example Project: A Blog



# Schema Validation

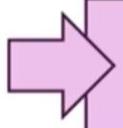


# Schema Validation

validationLevel

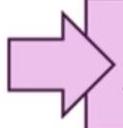
Which documents get validated?

strict



All inserts & updates

moderate



All inserts & updates  
to correct documents

validationAction

What happens if validation fails?

error



Throw error and deny  
insert/ update

warn



Log warning but  
proceed

# Data Modelling & Structuring – Things to Consider

In which Format will you fetch your Data?

How often will you fetch and change your Data?

How much data will you save (and how big is it)?

How is your Data related?

Will Duplicates hurt you (=> many Updates)?

Will you hit Data/ Storage Limits?

## Modelling Schemas

- Schemas should be modelled based on your application needs
- Important factors are: Read and write frequency, relations, amount (and size) of data

## Schema Validation

- You can define rules to validate inserts and update before writing to the database
- Choose your validation level and action based on your application requirements

## Modelling Relations

- Two options: Embedded documents or references
- Use embedded documents if you got one-to-one or one-to-many relationships and no app or data size reason to split
- Use references if data amount/ size or application needs require it or for many-to-many relations
- Exceptions are always possible => Keep your app requirements in mind!

# Working with Shell & Server

---

Beyond Start & Stop

# What's Inside This Module?

Start MongoDB Server as Process & Service

Configuring Database & Log Path (and Mode)

Fixing Issues

# MongoDB Compass

---

Exploring the Data Visually

# Diving Deeper Into **CREATE**

---

A Closer Look at Creating & Importing Documents

# What's Inside This Module?

Document Creation Methods (**CREATE**)

Importing Documents

# CREATE Documents

insertOne()

```
db.collectionName.insertOne({field: "value"})
```

insertMany()

```
db.collectionName.insertMany([
    {field: "value"}, 
    {field: "value"}])
```

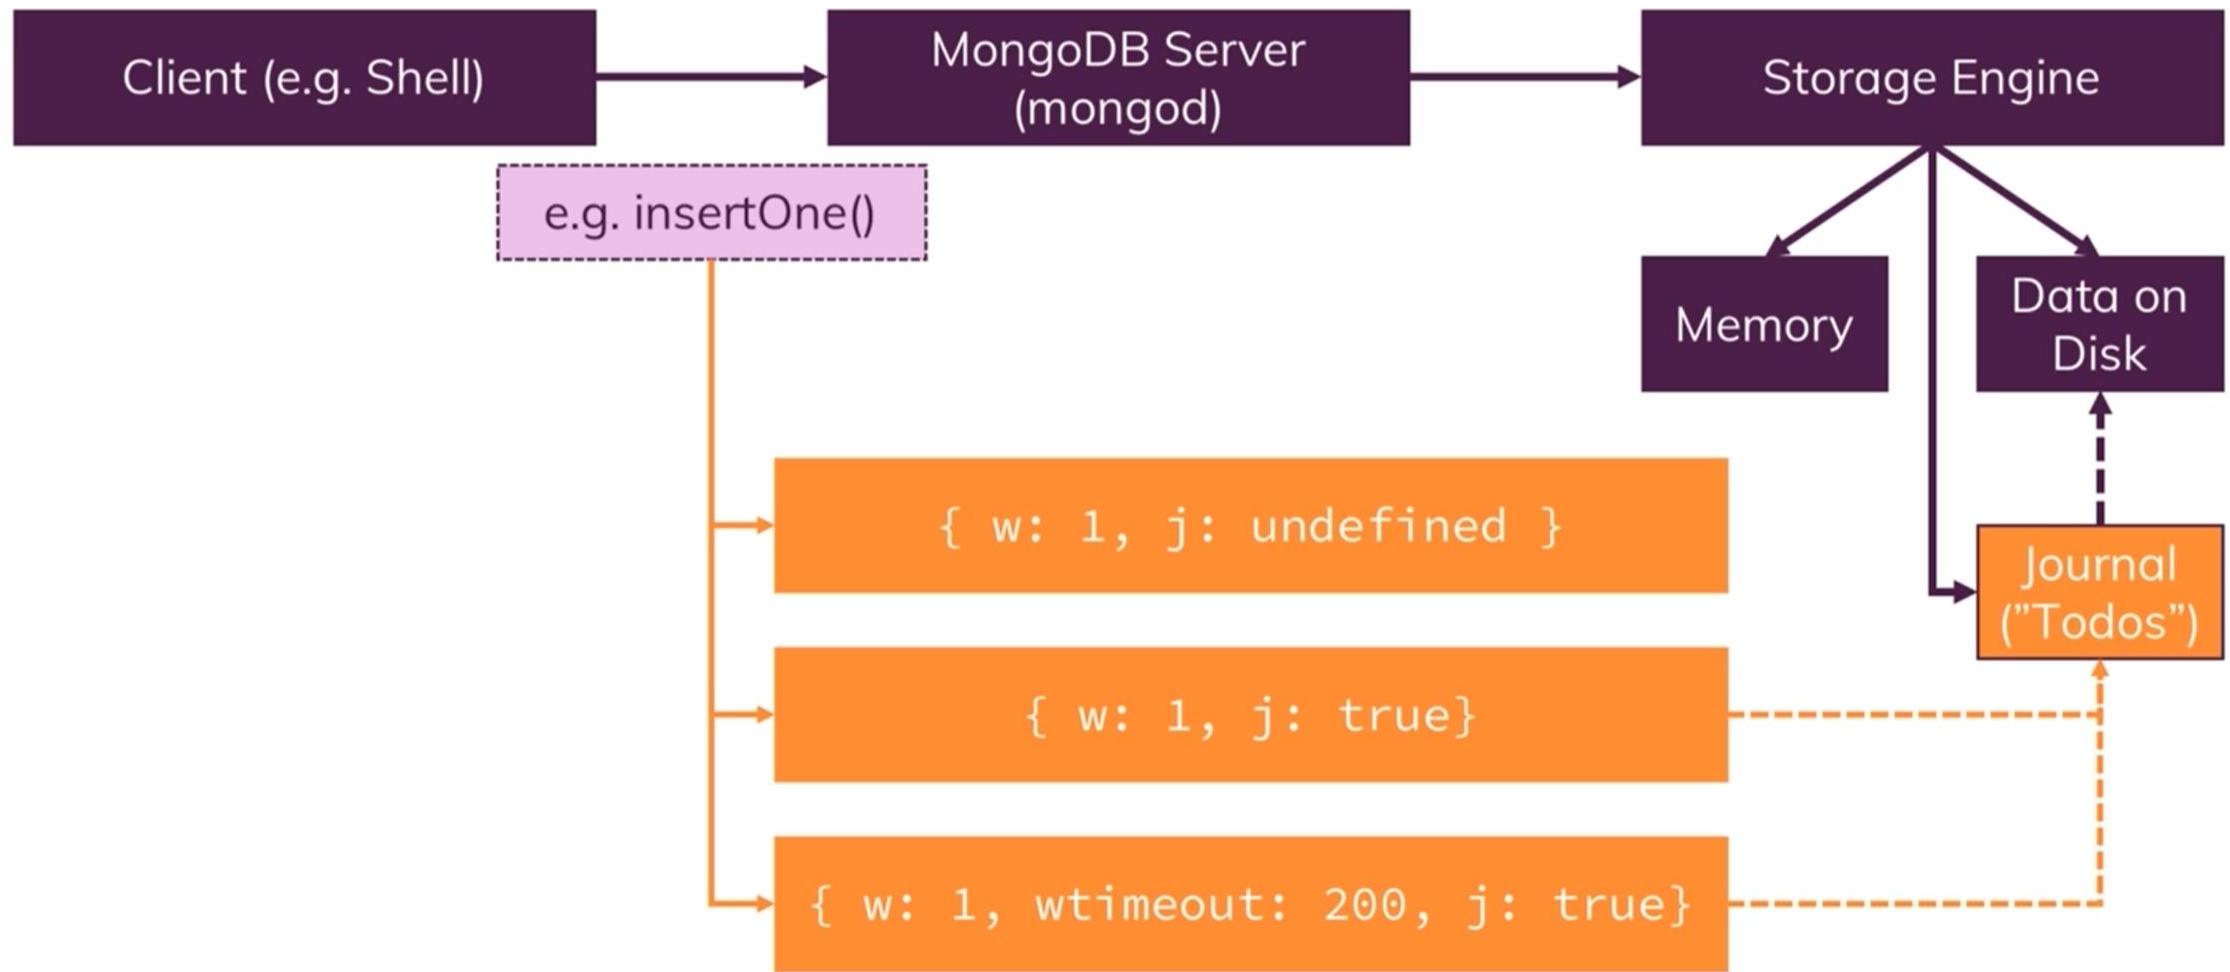
insert()

```
db.collectionName.insert()
```

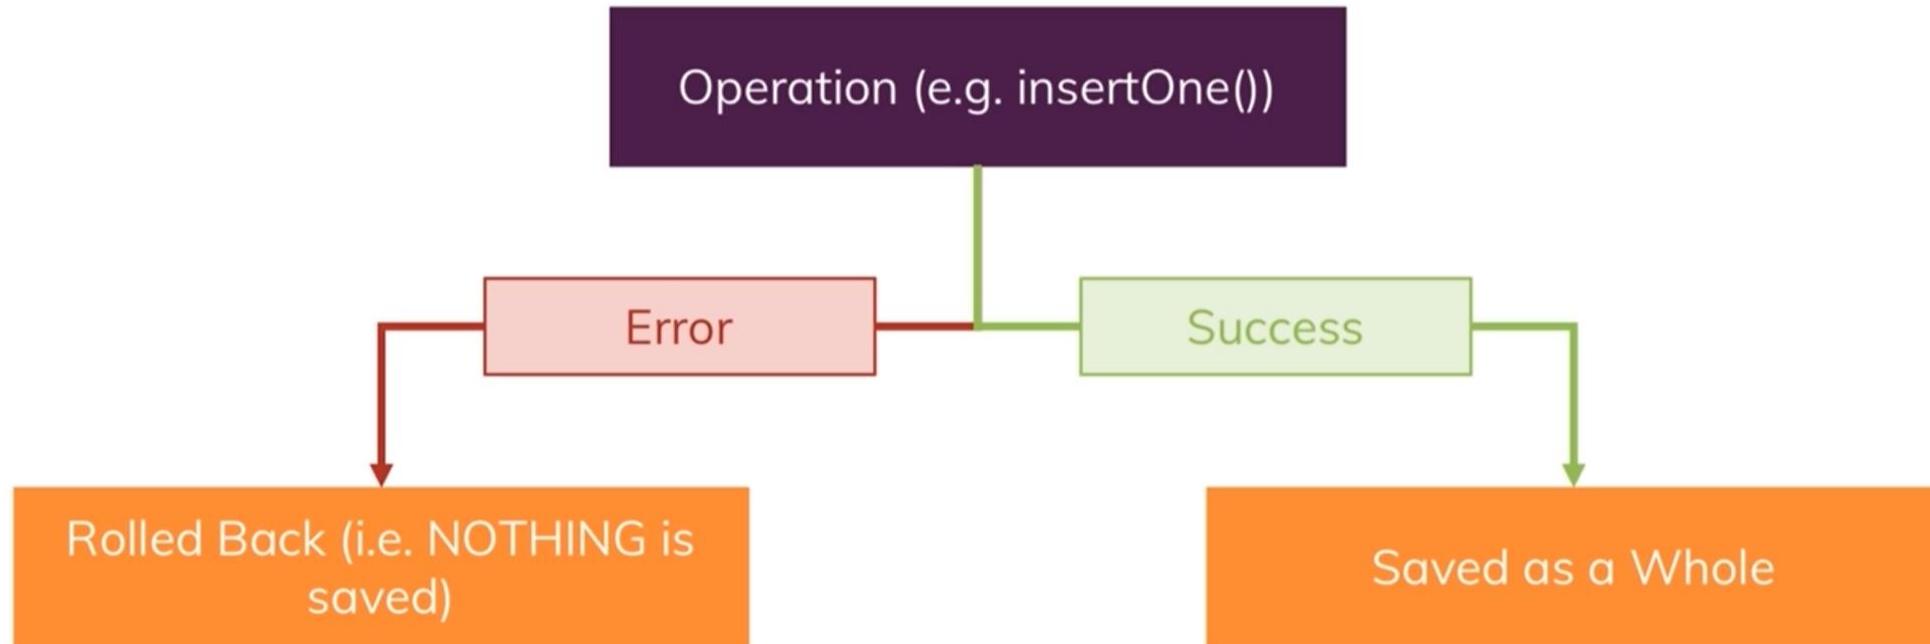
mongoimport

```
mongoimport -d cars -c carsList --drop --jsonArray
```

# WriteConcern



# What is “Atomicity”?



MongoDB CRUD Operations are Atomic on the Document Level (including Embedded Documents)

# Module Summary

## insertOne(), insertMany()

- You can insert documents with `insertOne()` (one document at a time) or `insertMany()` (multiple documents)
- `insert()` also exists but it's not recommended to use it anymore – it also doesn't return the inserted ids

## WriteConcern

- Data should be stored and you can control the “level of guarantee” of that to happen with the `writeConcern` option
- Choose the option value based on your app requirements

## Ordered Inserts

- By default, when using `insertMany()`, inserts are ordered – that means, that the inserting process stops if an error occurs
- You can change this by switching to “unordered inserts” – your inserting process will then continue, even if errors occurred
- In both cases, no successful inserts (before the error) will be rolled back

# **READiNg DOCuMEnTS wITH OPeRAToRS**

---

Accessing the Required Data Efficiently

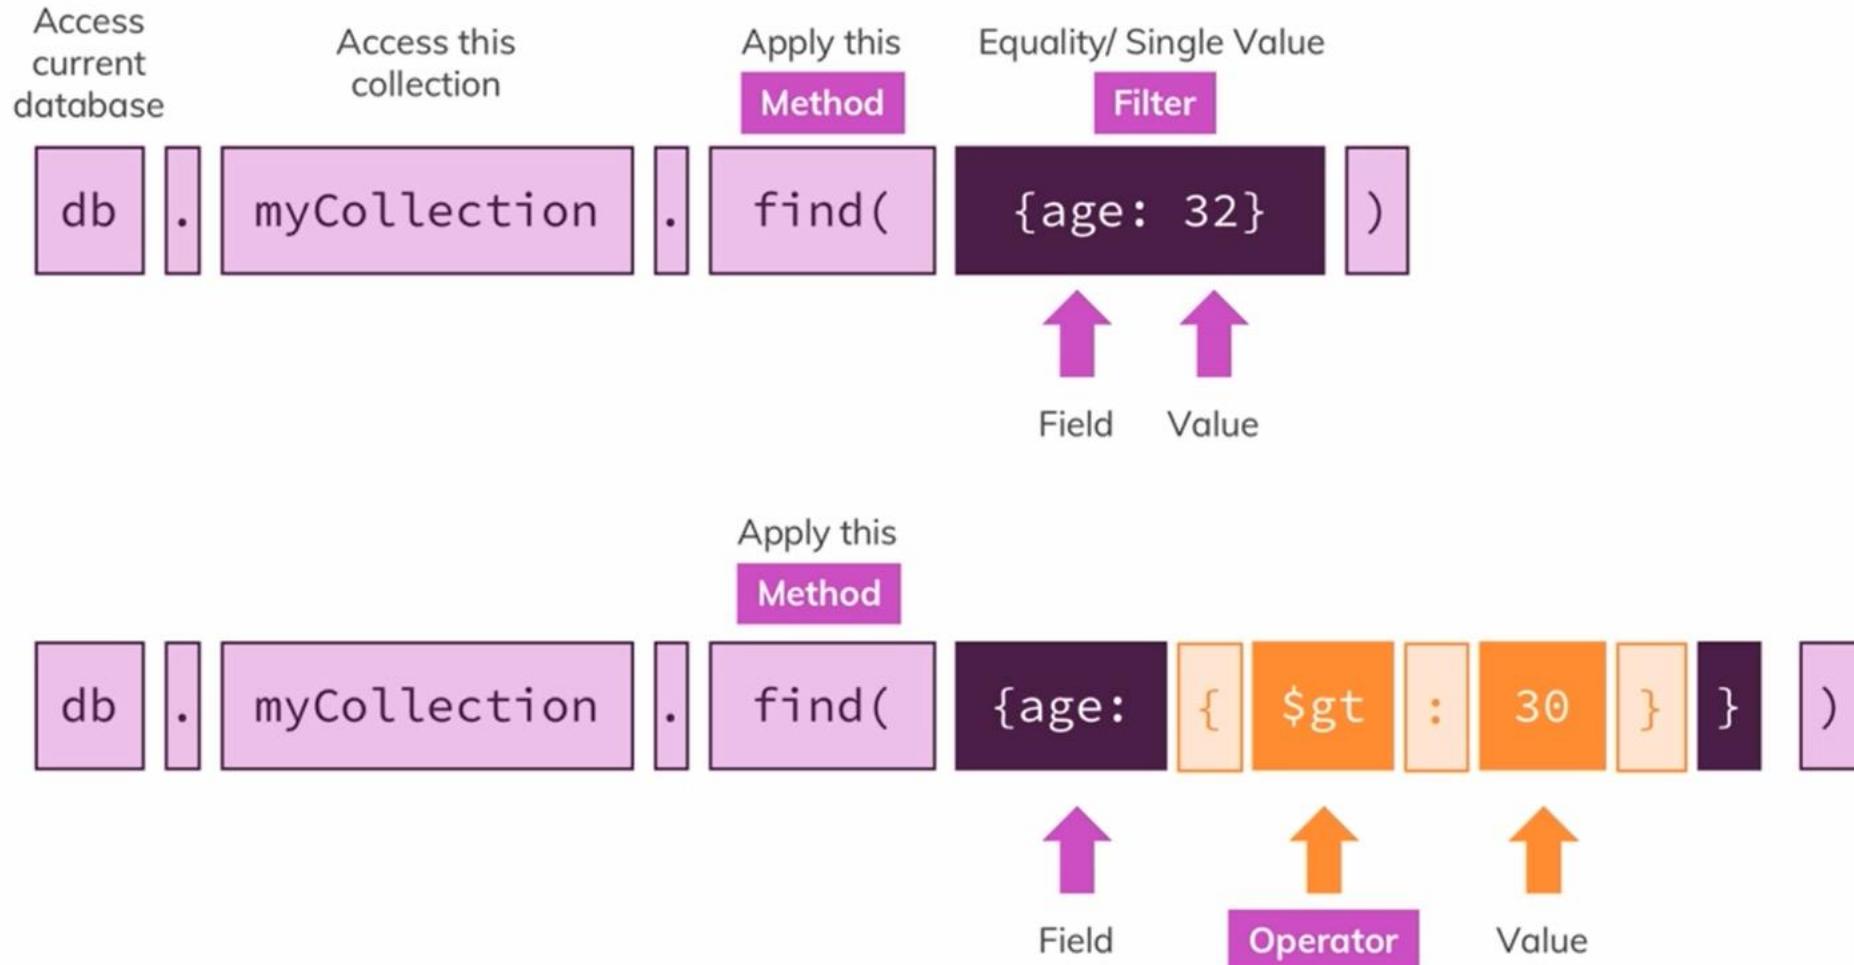
# What's Inside This Module?

Methods, Filters & Operators

Query Selectors (**READ**)

Projection Operators (**READ**)

# Methods, Filters & Operators



# Operators

Read

Update

Query & Projection

Update

Query Modifiers

Aggregation

Query Selectors

Fields

Pipeline Stages

Projection Operators

Arrays

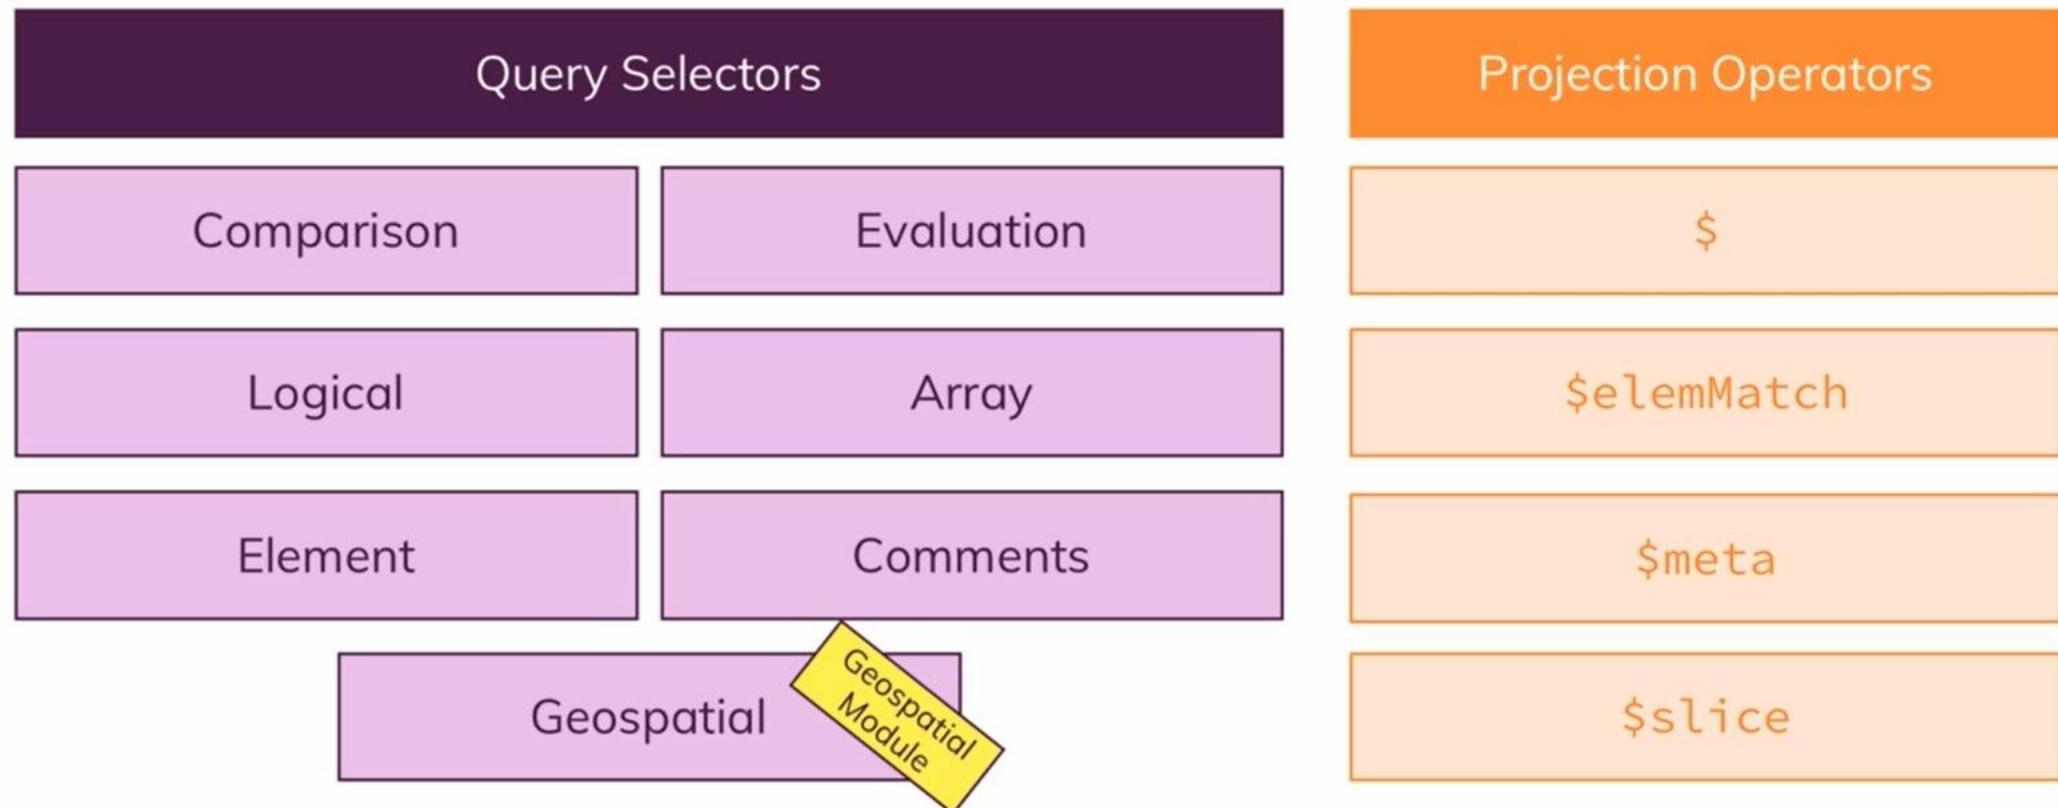
Change Stream  
Deprecated

Aggregation  
Module  
Operators

# How Operators Impact our Data

Type	Purpose	Changes Data?	Example
Query Operator	Locate Data	🚫	\$eq
Projection Operator	Modify data presentation	🚫	\$
Update Operator	Modify + add additional data	✓	\$inc

# Query Selectors & Projection Operators



# Indexes

---

Retrieving Data Efficiently

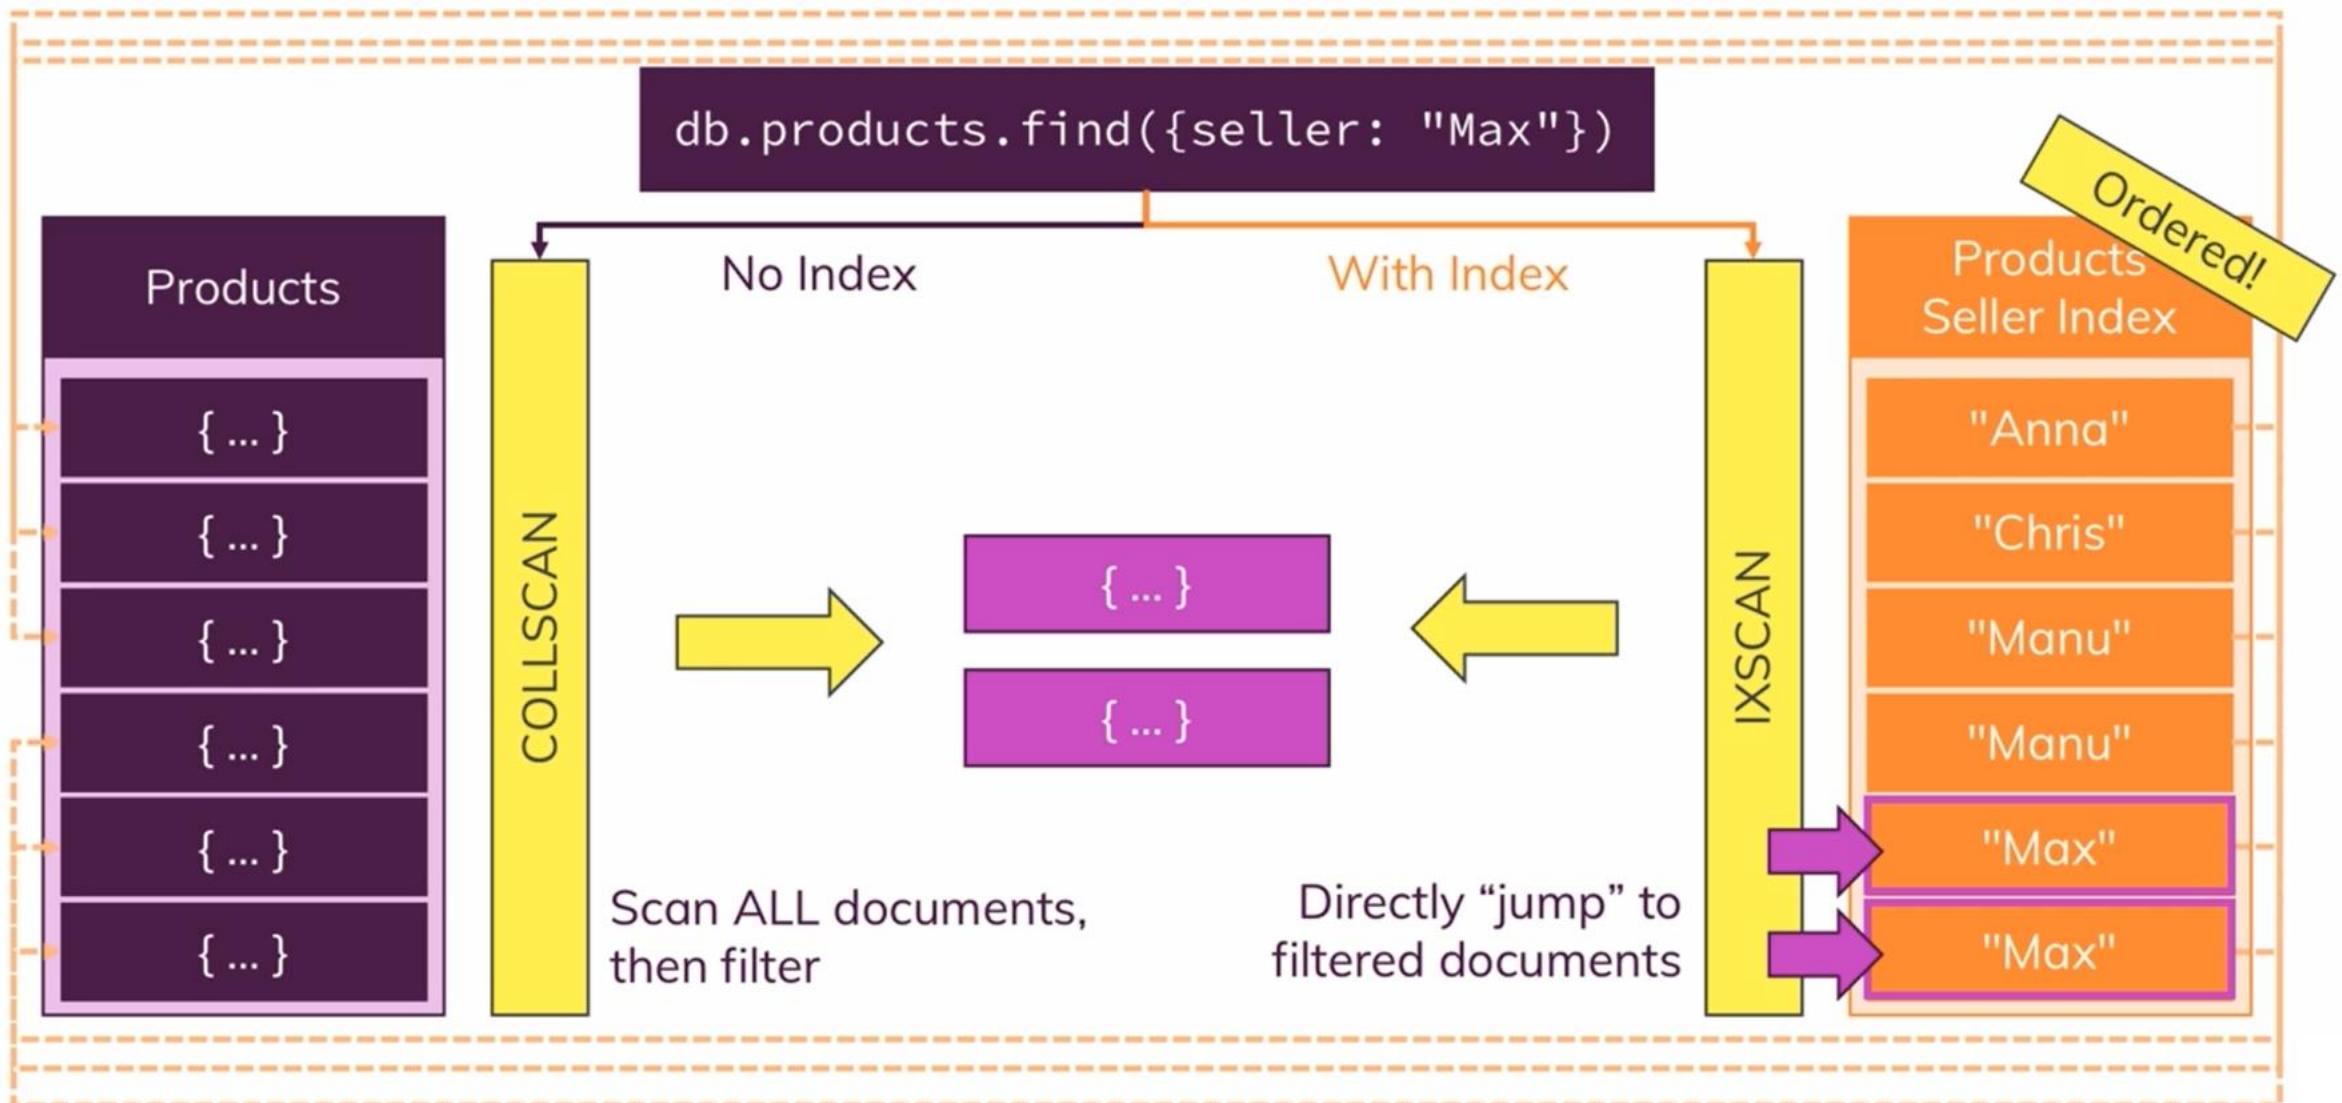
# What's Inside This Module?

What are Indexes?

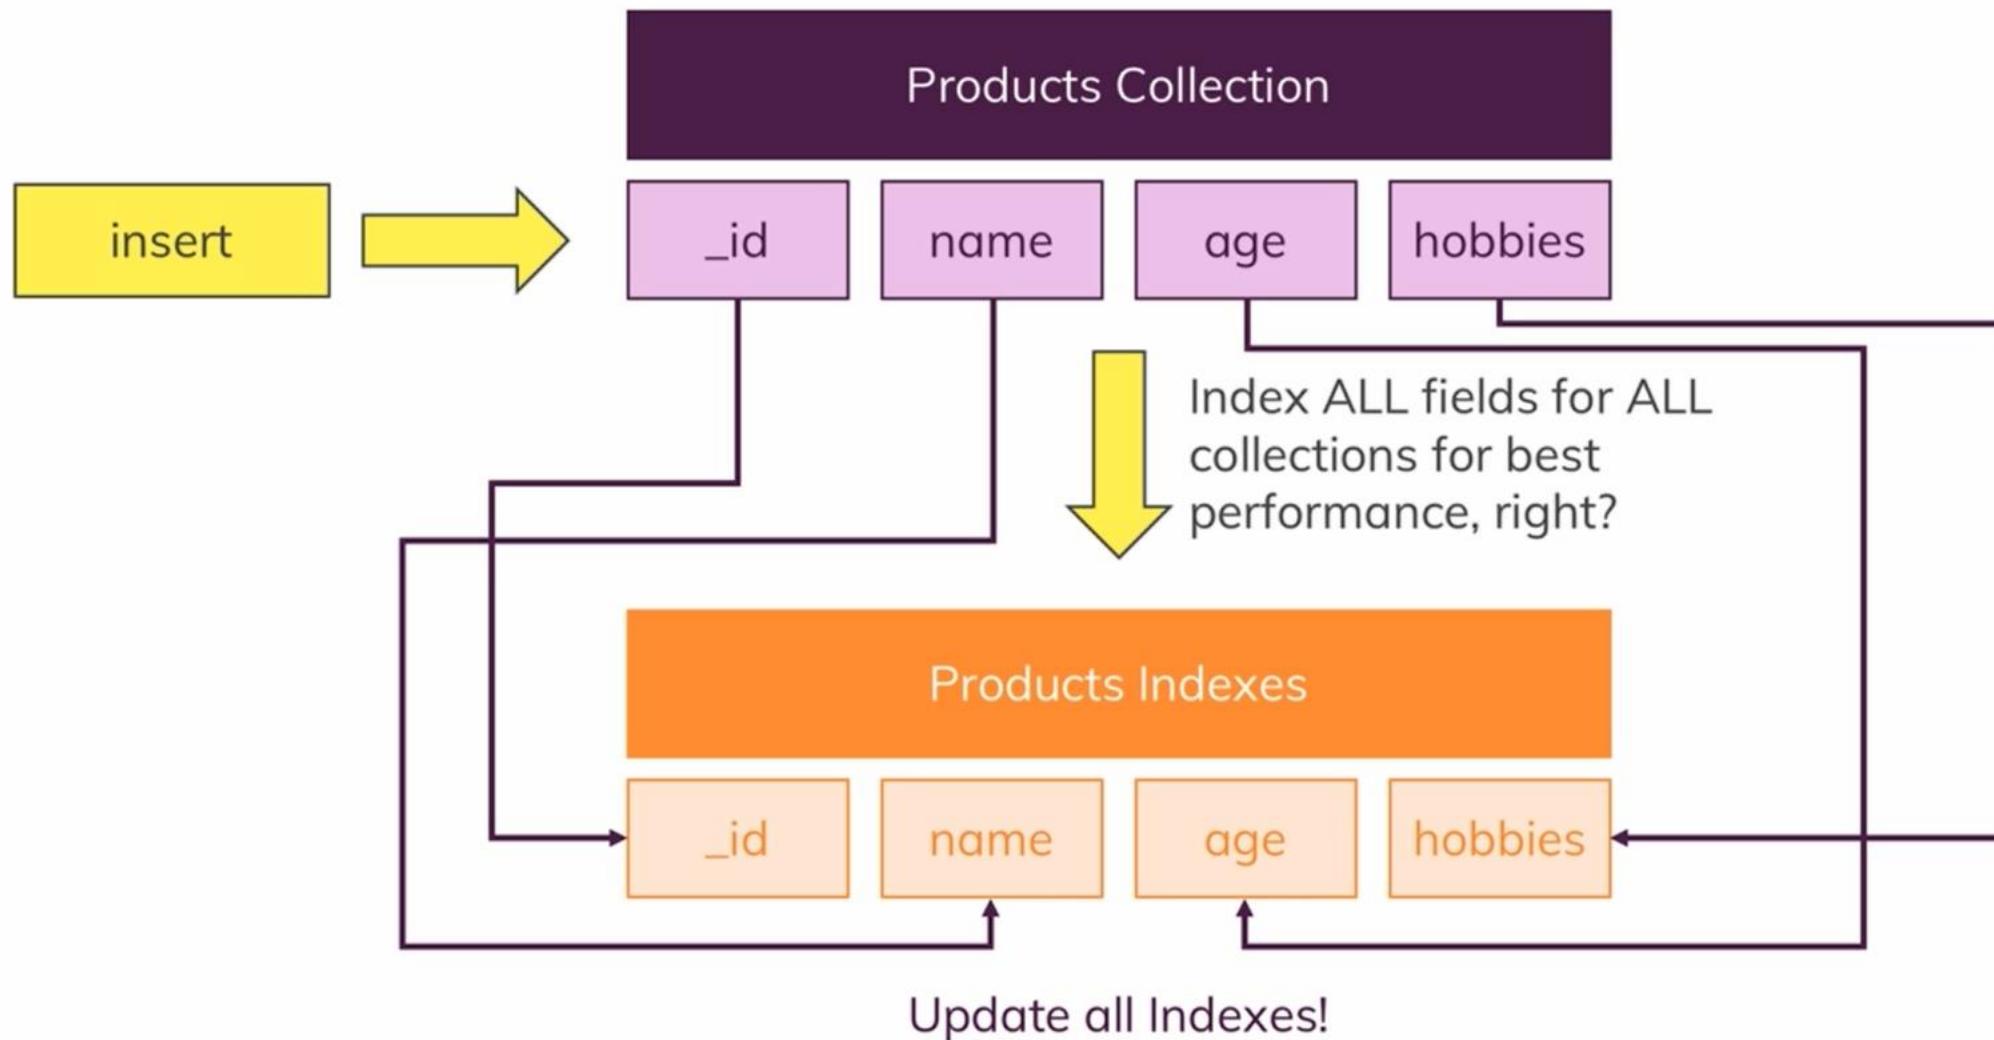
Different Types of Indexes

Using & Optimizing Indexes

# Why Indexes?



# Don't Use Too Many Indexes!



# Geospatial Queries

---

Finding Places

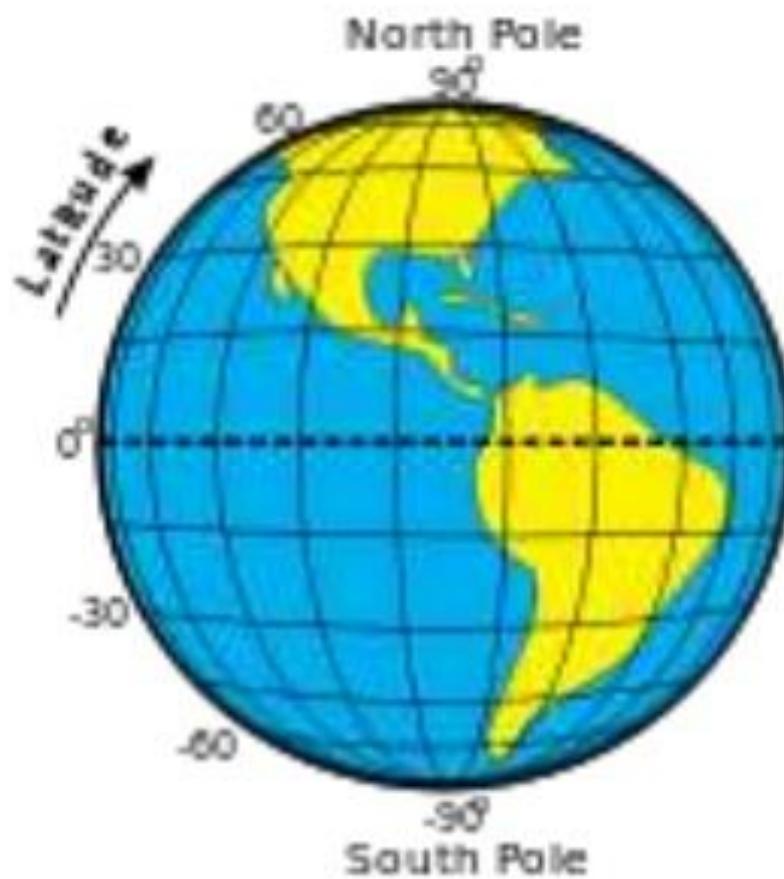
# What's Inside This Module?

Storing Geospatial Data in GeoJSON Format

Querying Geospatial Data

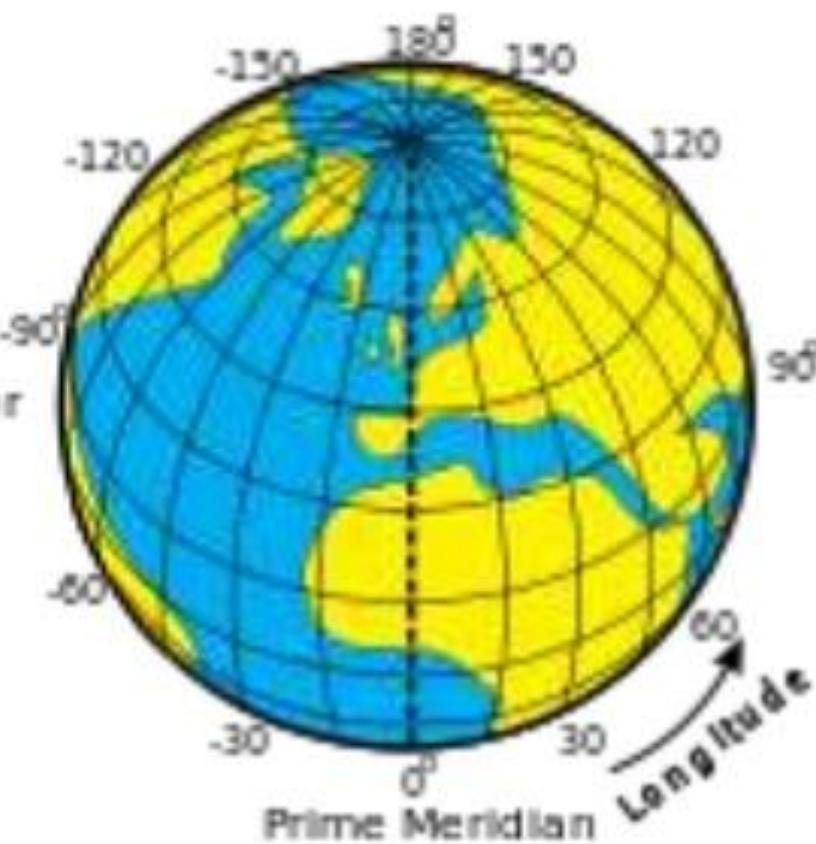
# Horizontal Lines

Latitude | **Parallels**



# Vertical Lines

Longitude | **Meridians**

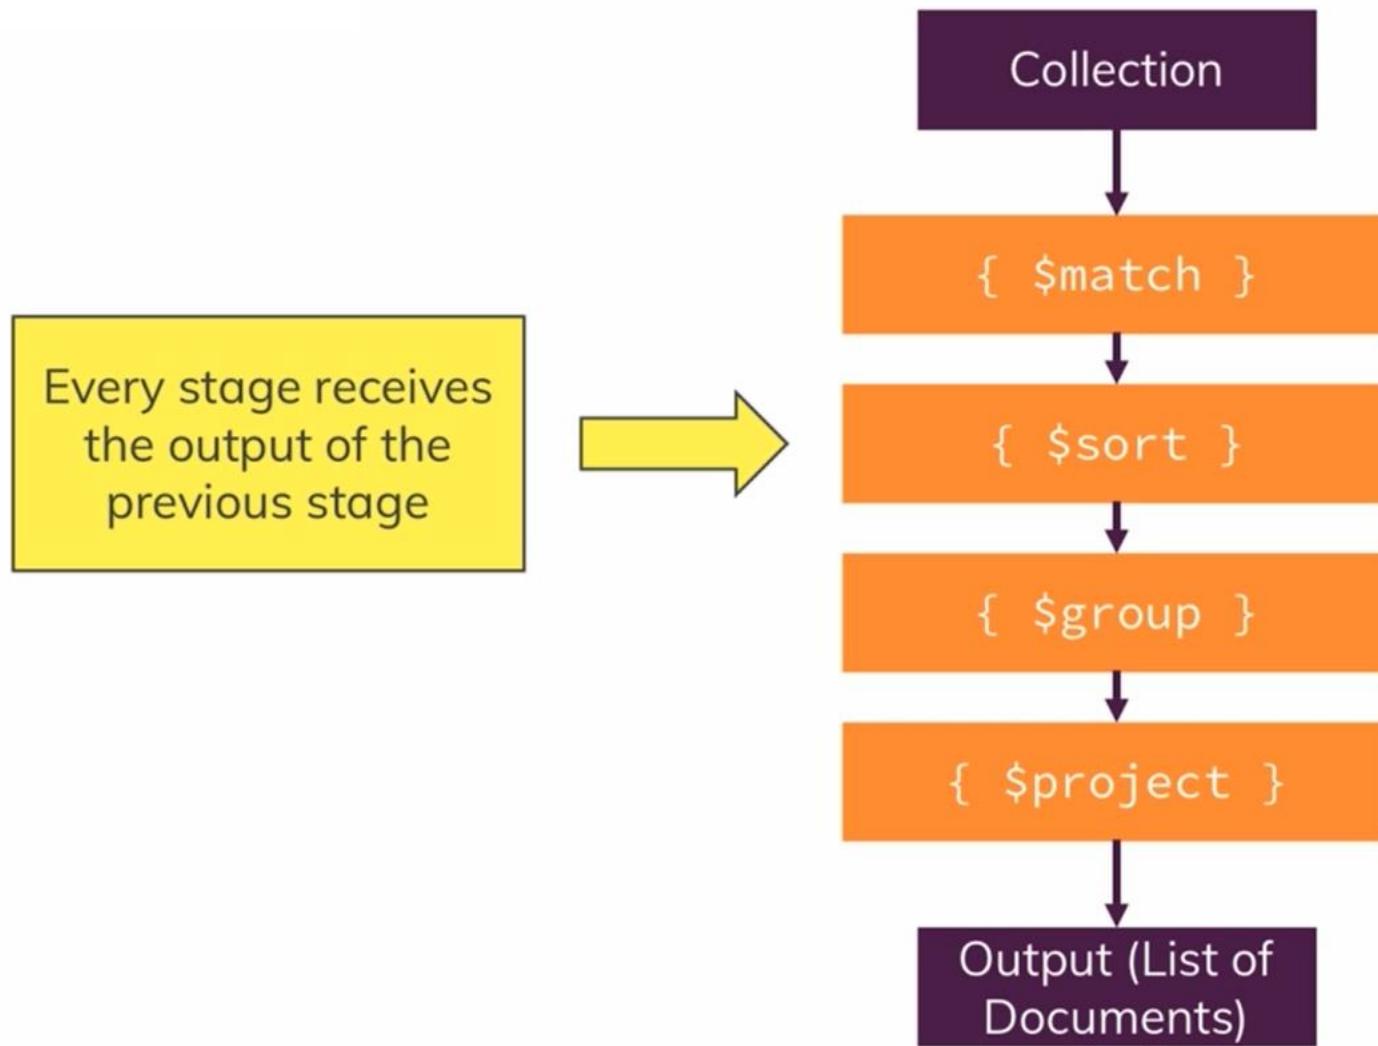


# Using the Aggregation Framework

---

Retrieving Data Efficiently & In a Structured Way

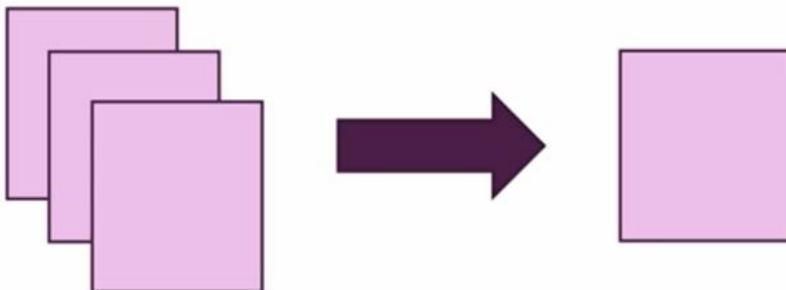
# What is the “Aggregation Framework”?



# \$group vs \$project

\$group

n:1



Sum, Count, Average, Build Array

\$project

1:1



Include/ Exclude Fields, Transform  
Fields (within a Single Document)

## MONGO DB DATABASE

Self Managed/  
Enterprise

Cloud Manager /  
Ops Manager

Compass / GUI Tool

BI Connectors

MongoDB Charts

Atlas Cloud

Mobile

Stitch

Serverless Query API

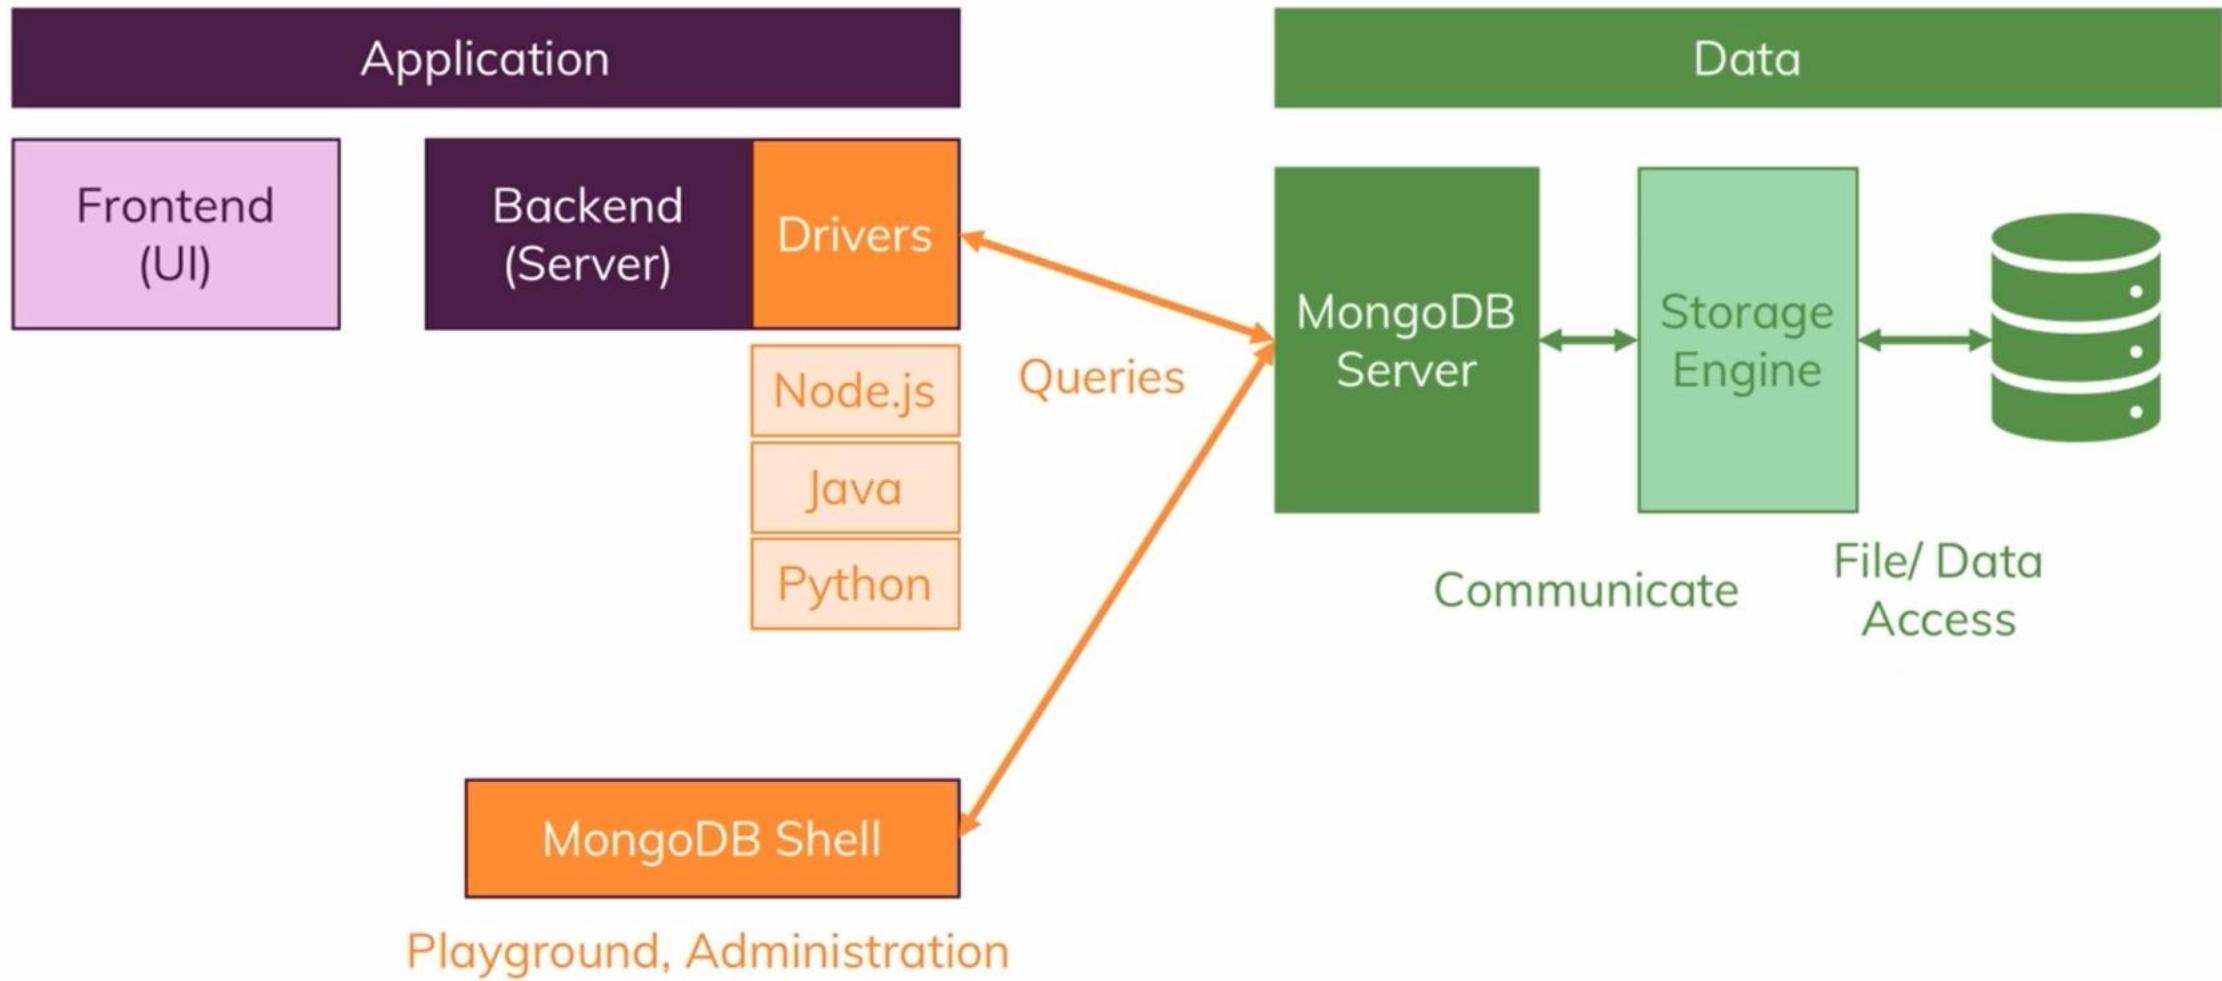
Serverless Functions

Database Triggers

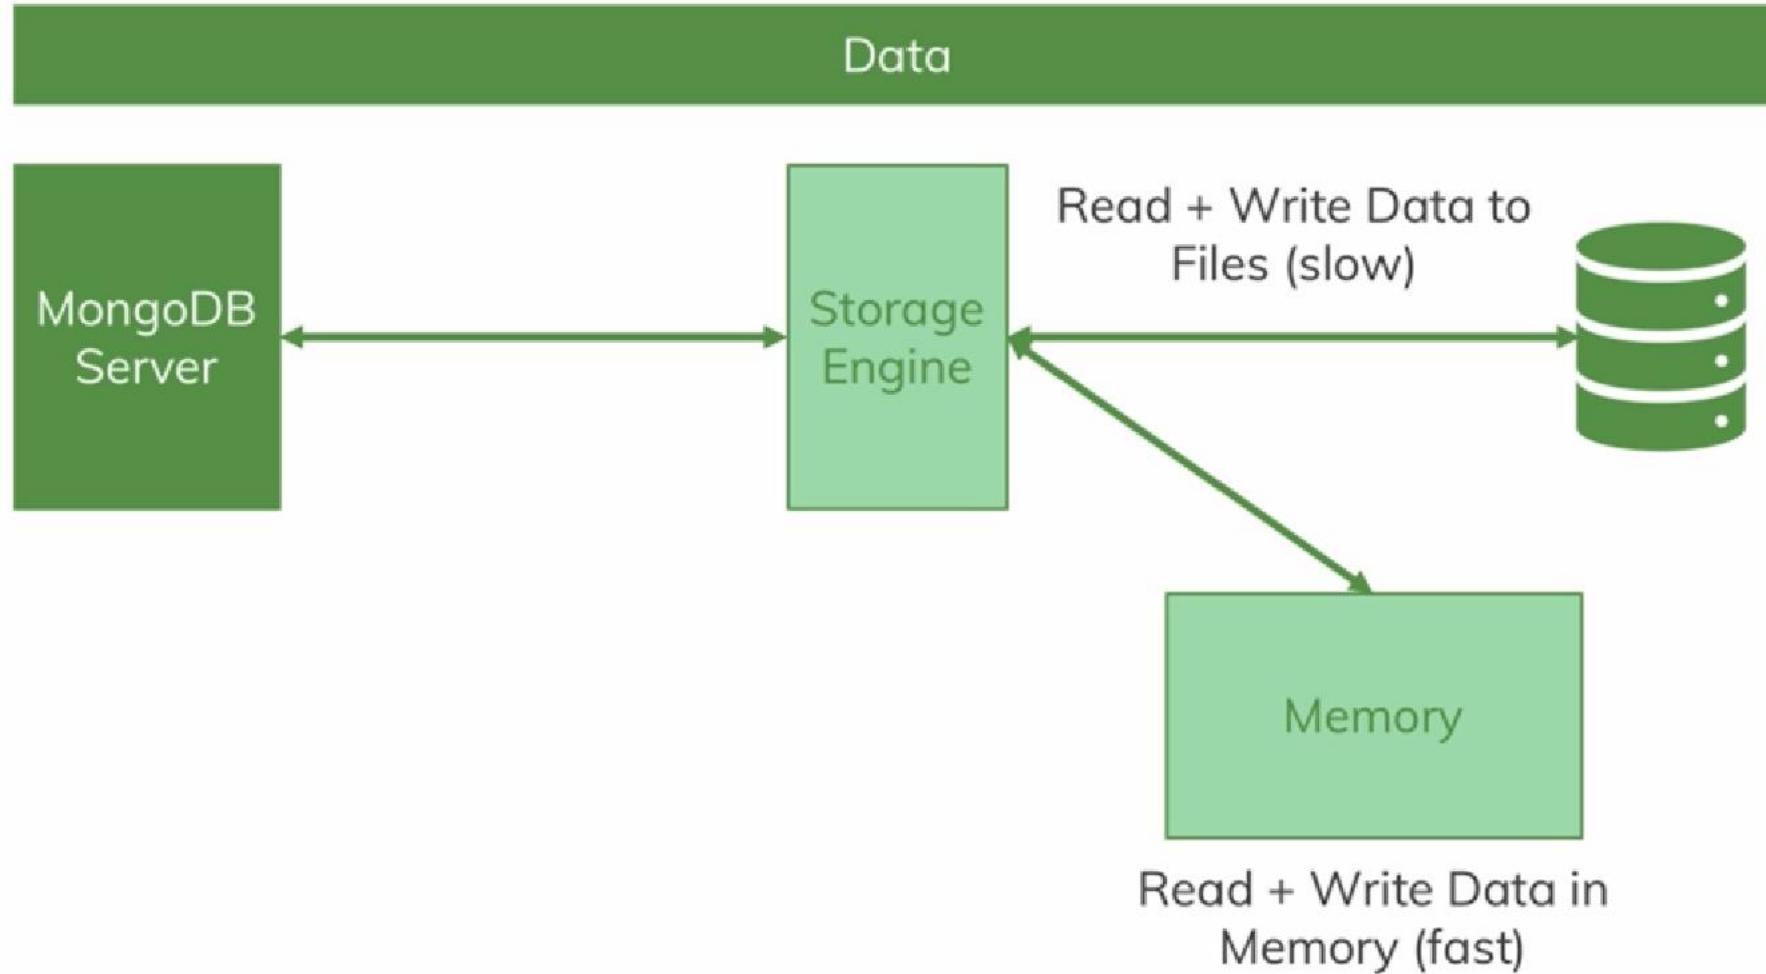
Real-Time Sync



# Working with MongoDB



# A Closer Look



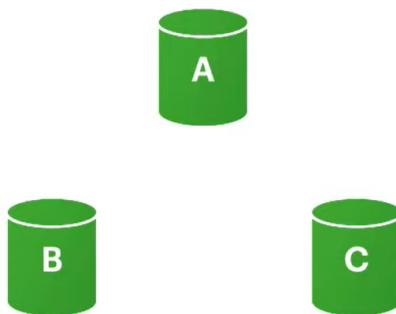
# Cluster

## Replica Sets

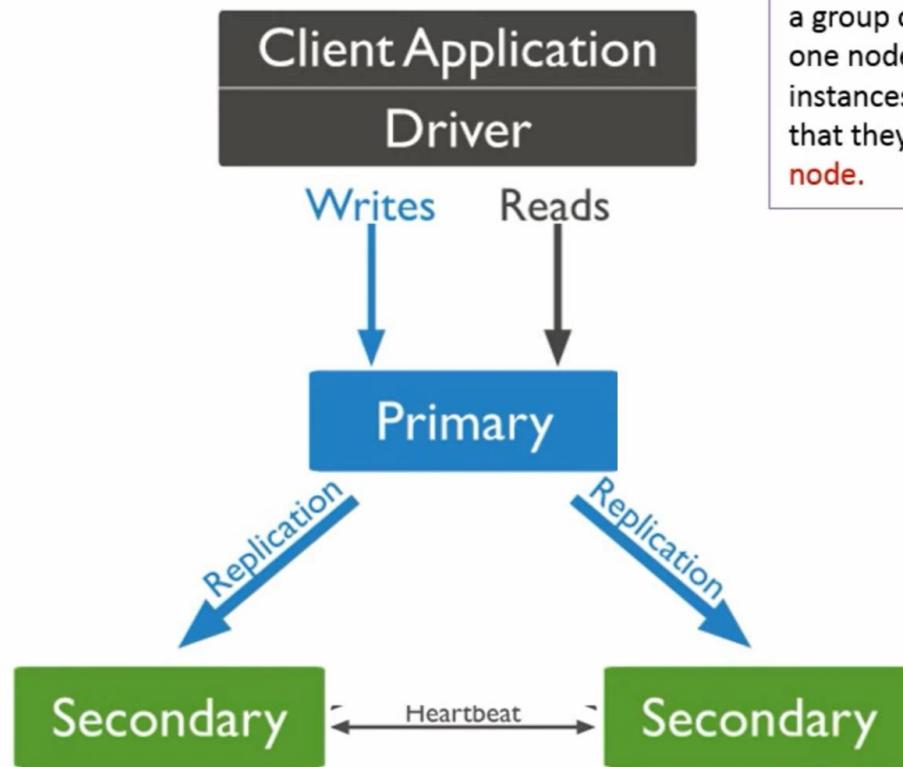
- ✓ Maintain same dataset
- ✓ Provides redundancy and high availability
- ✓ Load balancing
- ✓ Basis for all Production deployments

## Sharded Cluster

- ✓ Distributing Data between different machines
- ✓ Useful when u have very large datasets



# Replica Set



MongoDB achieves replication by the use of **replica set**. A **replica set** is a group of **mongod** instances that host the same data set. In a replica, one node is primary node that receives all write operations. All other instances, such as secondaries, apply operations from the primary so that they have the same data set. **Replica set can have only one primary node.**

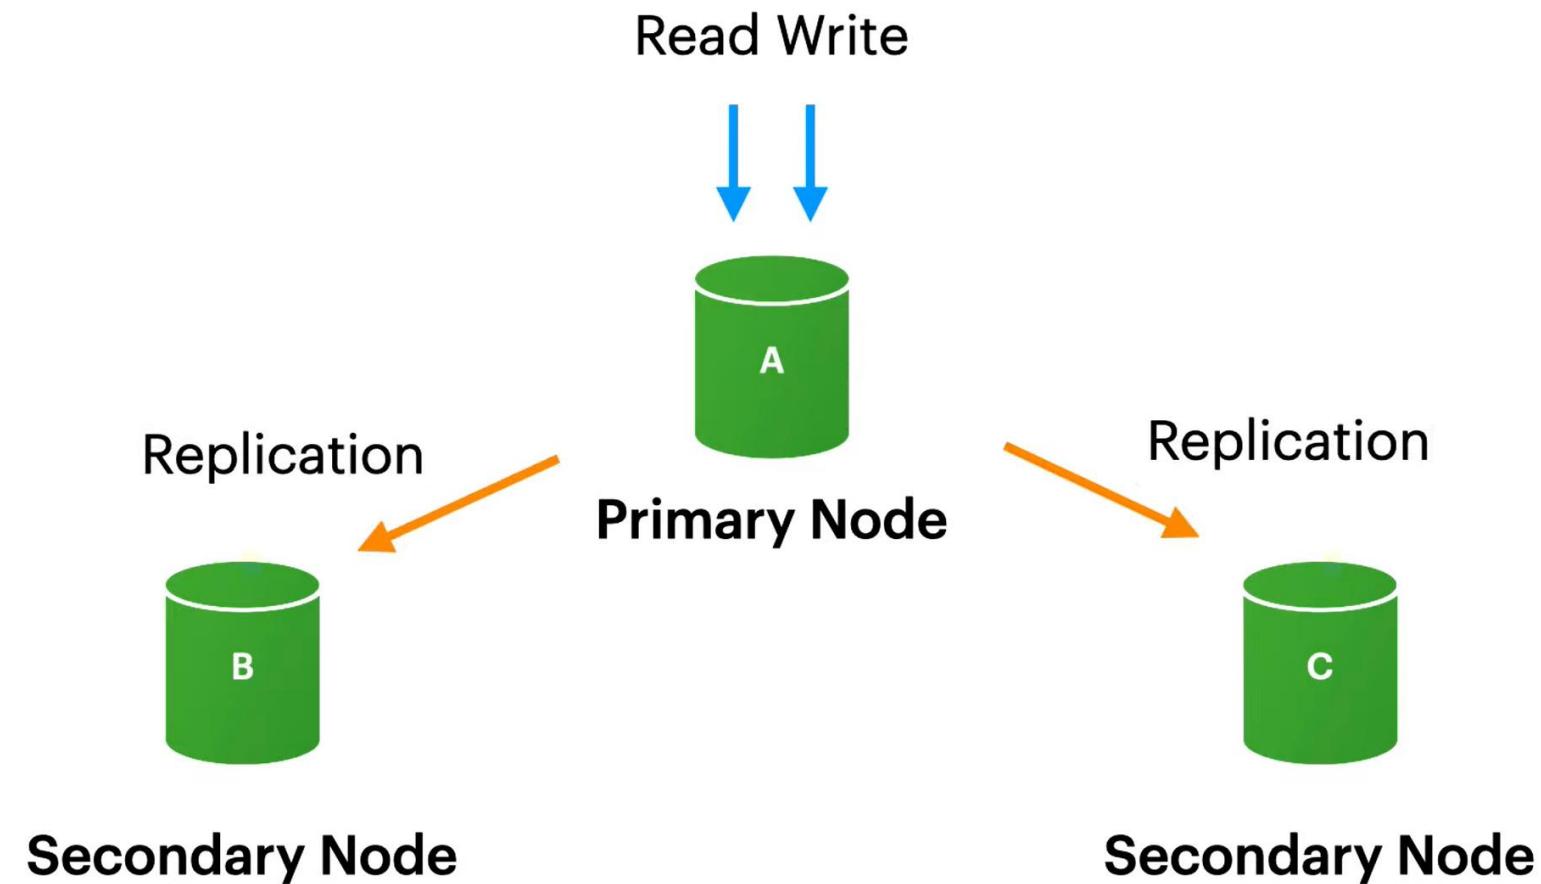
Single Node ReplicaSet

Primary

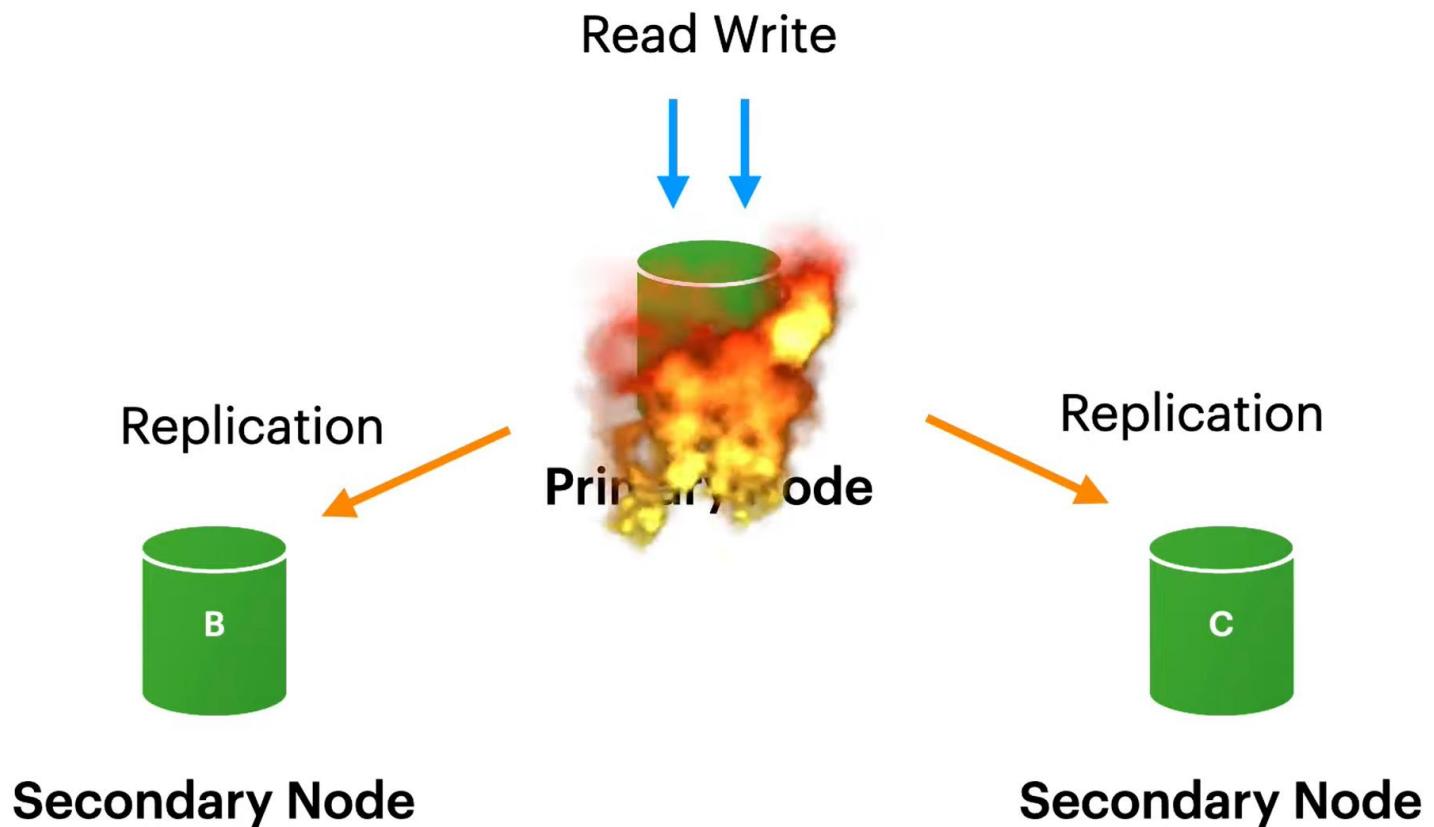
3 Nodes ReplicaSet

1. Primary
2. Secondary1
3. Secondary2

# Replica Set

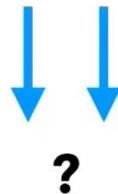


# Replica Set



# Replica Set

Read Write



Election takes place

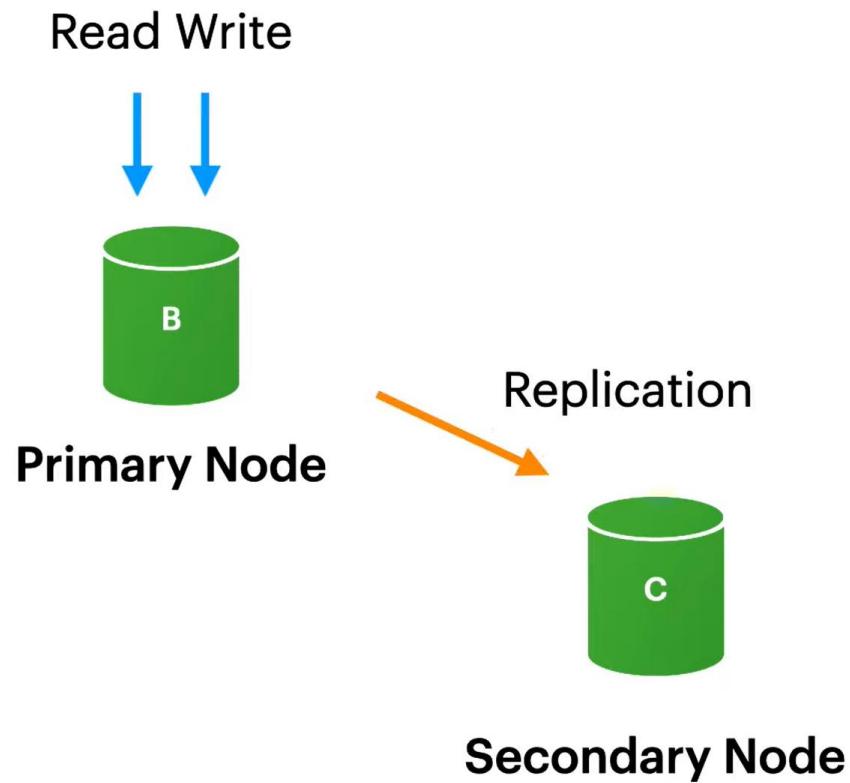


Secondary Node

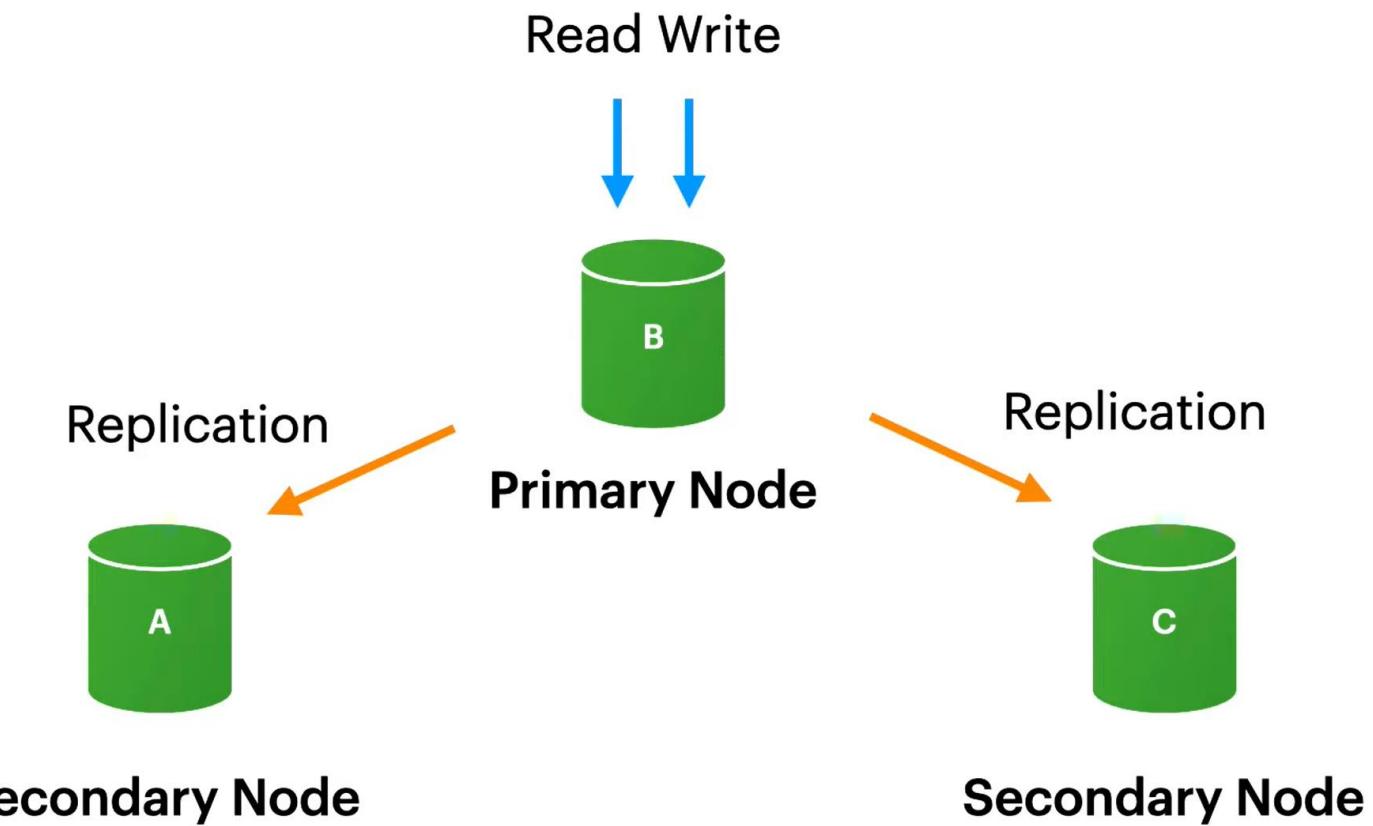


Secondary Node

# Replica Set

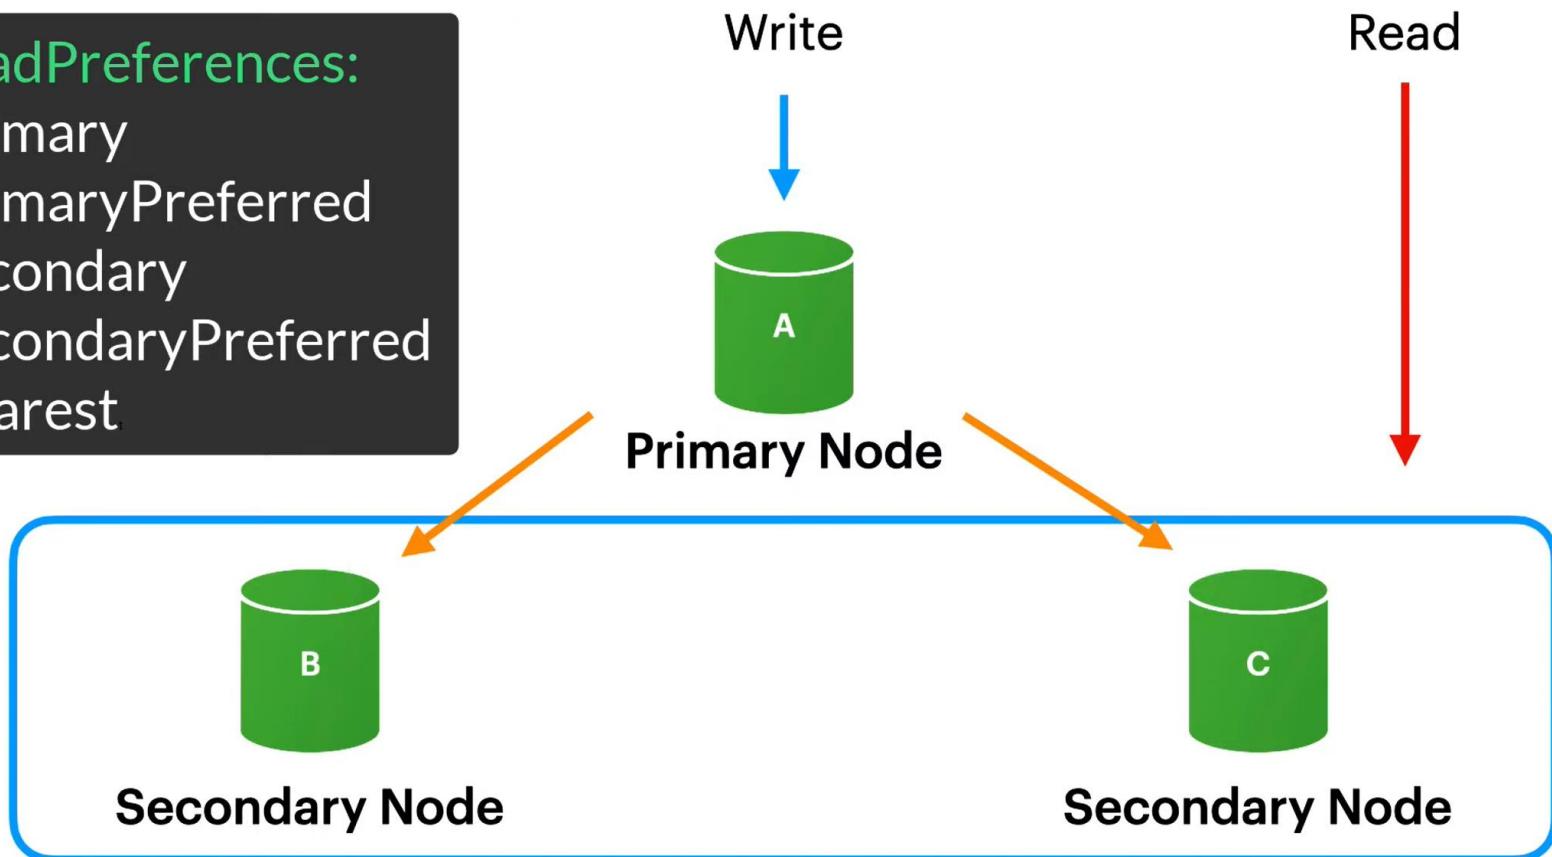


# Replica Set



# Load Balancing

`readPreferences:`  
primary  
primaryPreferred  
secondary  
secondaryPreferred  
nearest



# Cluster Setup

Host: localhost (127.0.0.1)



## Primary Node

- Port: 2717
- Db path: /mongos/db1
- `mongod --port 2717 --dbpath /mongos/db1 --replicaSet myReplicaSet`



## Secondary Node

- Port: 2727
- Db path: /mongos/db2
- `mongod --port 2727 --dbpath /mongos/db2 --replicaSet myReplicaSet`



## Secondary Node

- Port: 2737
- Db path: /mongos/db3
- `mongod --port 2737 --dbpath /mongos/db3 --replicaSet myReplicaSet`

# Commands

## On Primary Node

**rs.status()**

- Tells the status of replica set

**rs.initiate()**

- Initiates a replica set on primary node

**rs.add("localhost:2727")**

- Adds a member to replica set

**rs.remove("localhost:2727")**

- Removes a member from replica set

# Connecting to ReplicaSet

## MongoDB URI

typically

- `mongodb://localhost:27017`

For, Replica sets

- `mongodb://localhost:2717,localhost:2727,localhost:2737/?replicaSet=myRepl`

# THANK YOU

Corporate Trainer & Motivational Speaker  
9035351965

